Azure Databricks Architecture and Security

databricks

Quick Agenda

- Azure Databricks Platform Architecture
- Azure Databricks Security
  - Data Protection
  - IAM/Auth
  - Network Security
  - Compliance

Azure Databricks Platform Architecture

Azure Databricks

AZURE DATA SOURCES

Blob Storage

Data Lake Store

SQL Data Warehouse

Cosmos DB

Event Hub

IoT Hub

Azure Data Factory

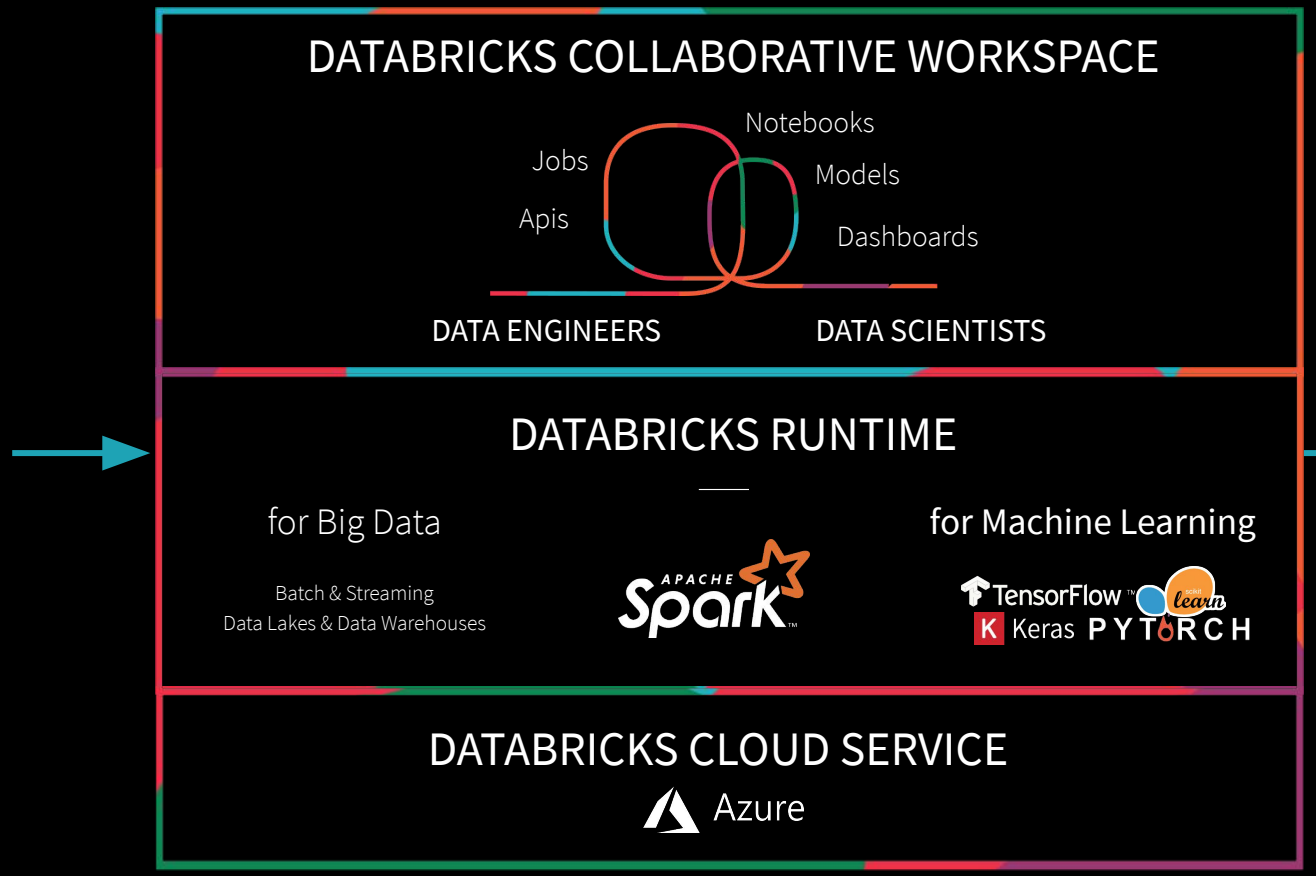

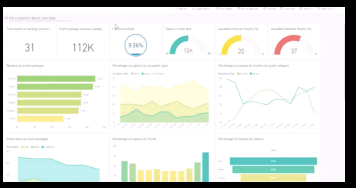

Power BI

BI Reporting
Dashboards

Microsoft Active Directory

Security Integration

Azure Portal
One-Click setup
Unified Billing

databricks

Azure Databricks Platform Architecture

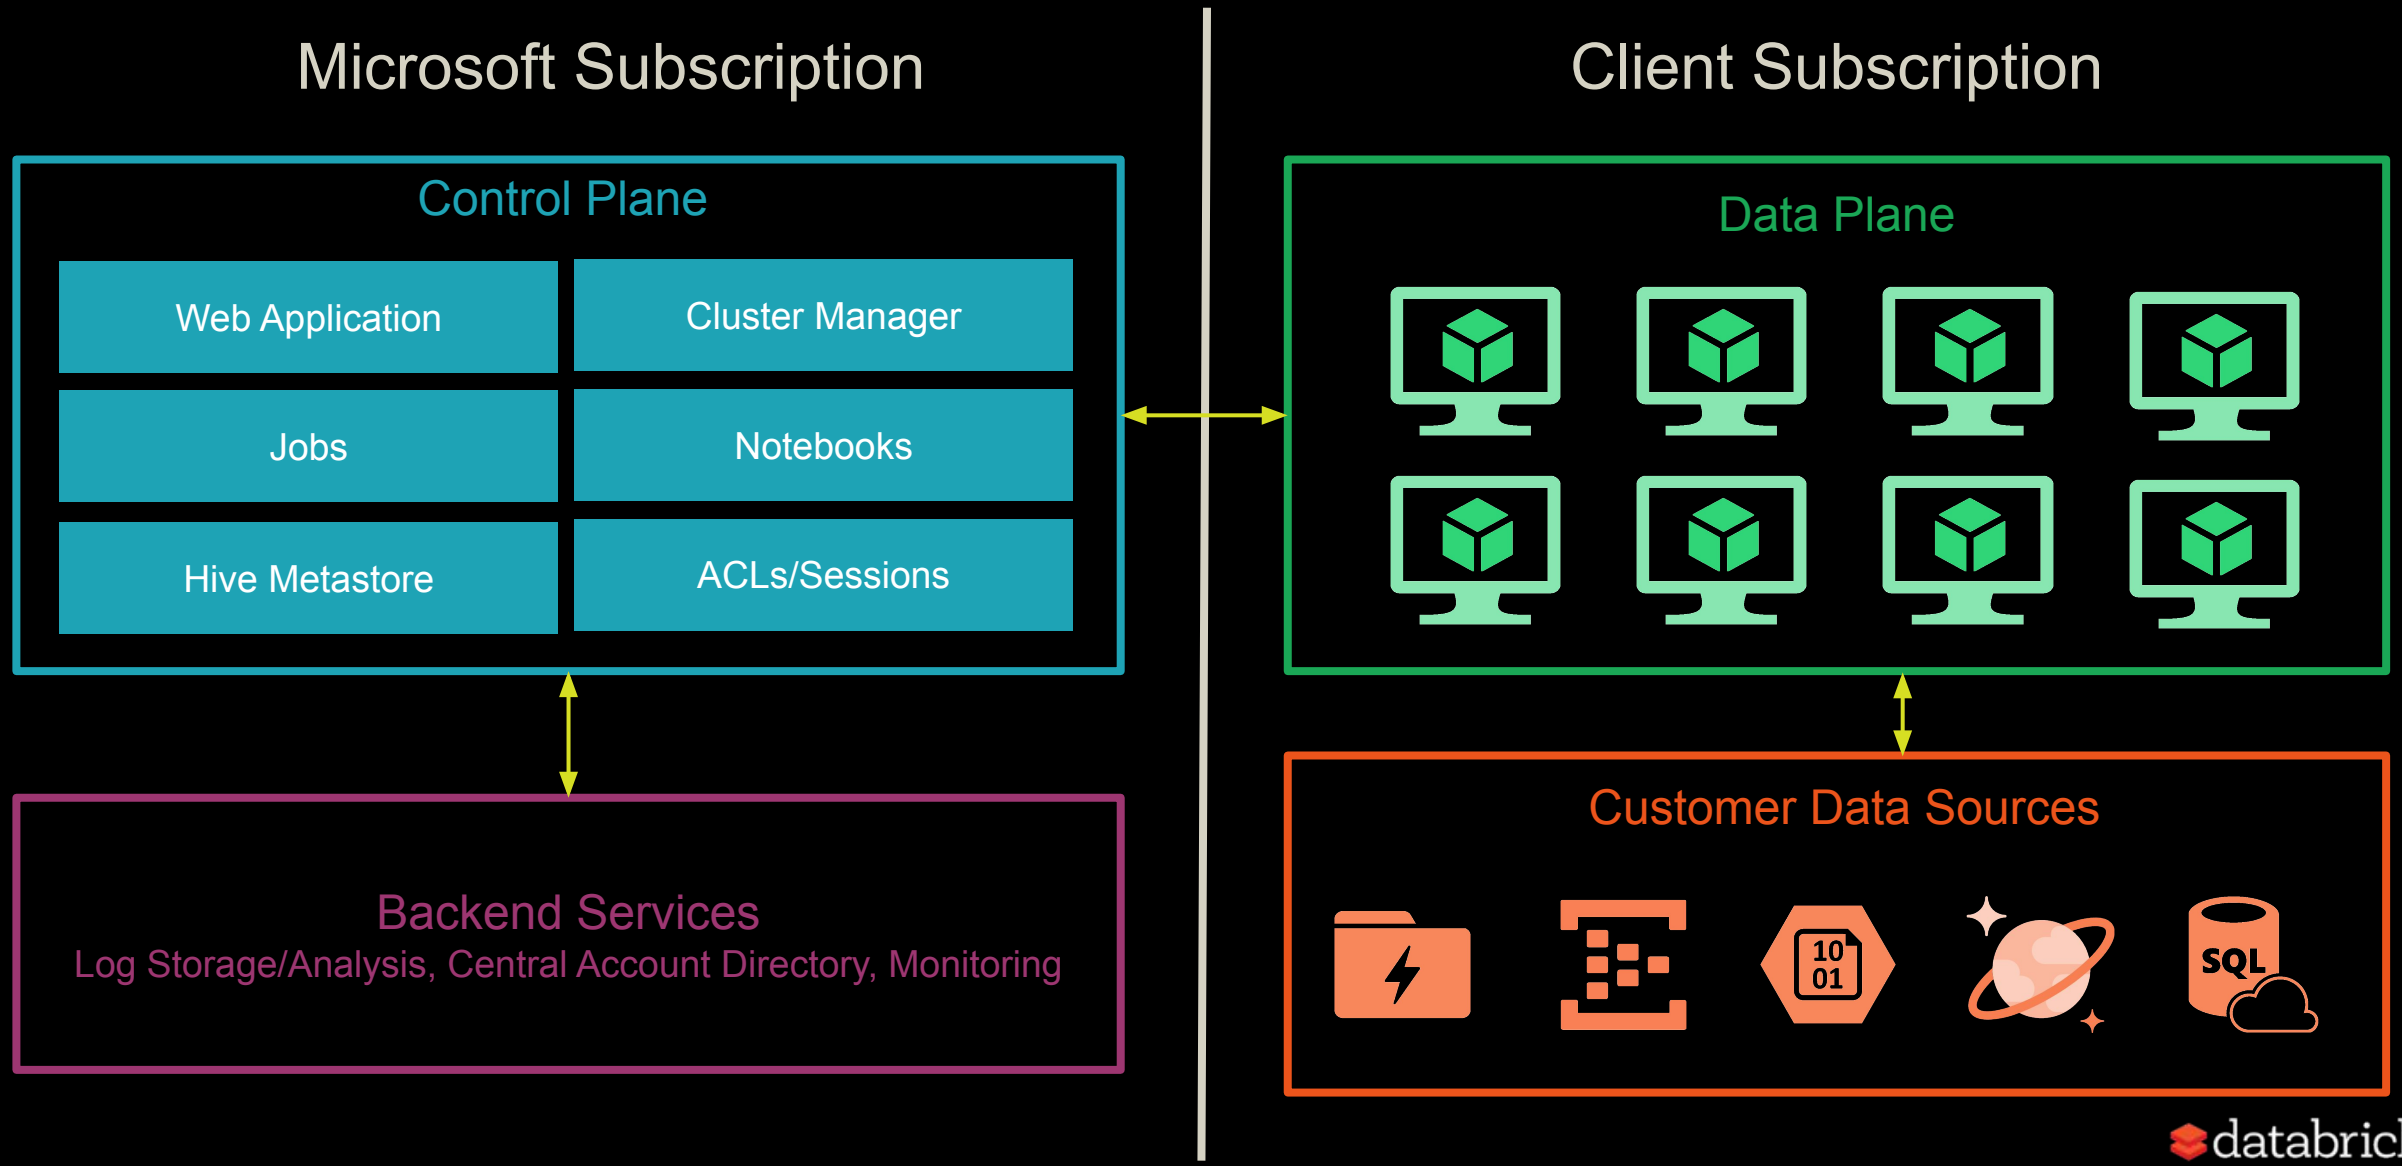

databricks

Azure Databricks Platform Architecture Cont'd

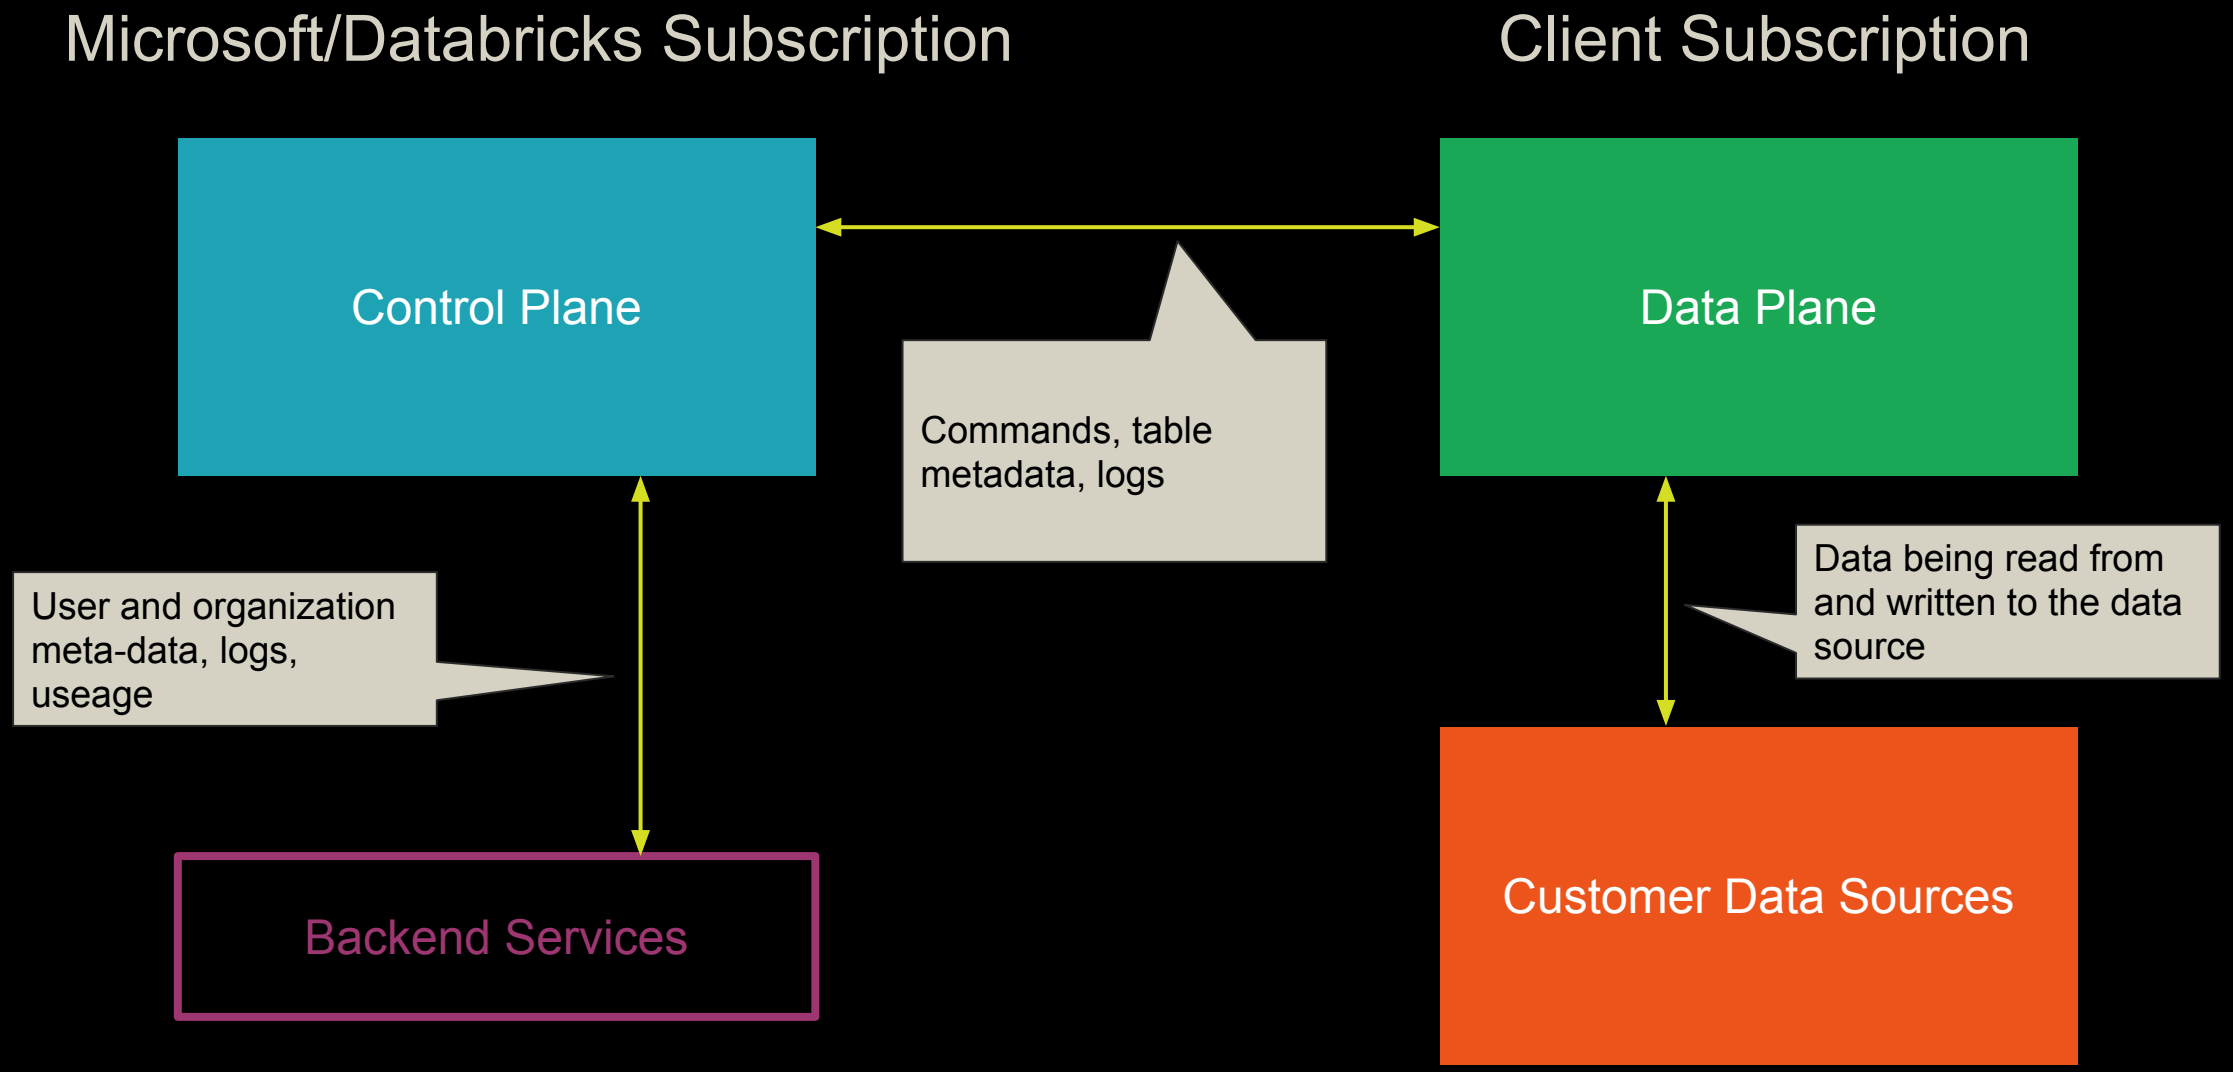

databricks

Azure Databricks Platform Architecture

Standard Deployment View with no inter-node TLS

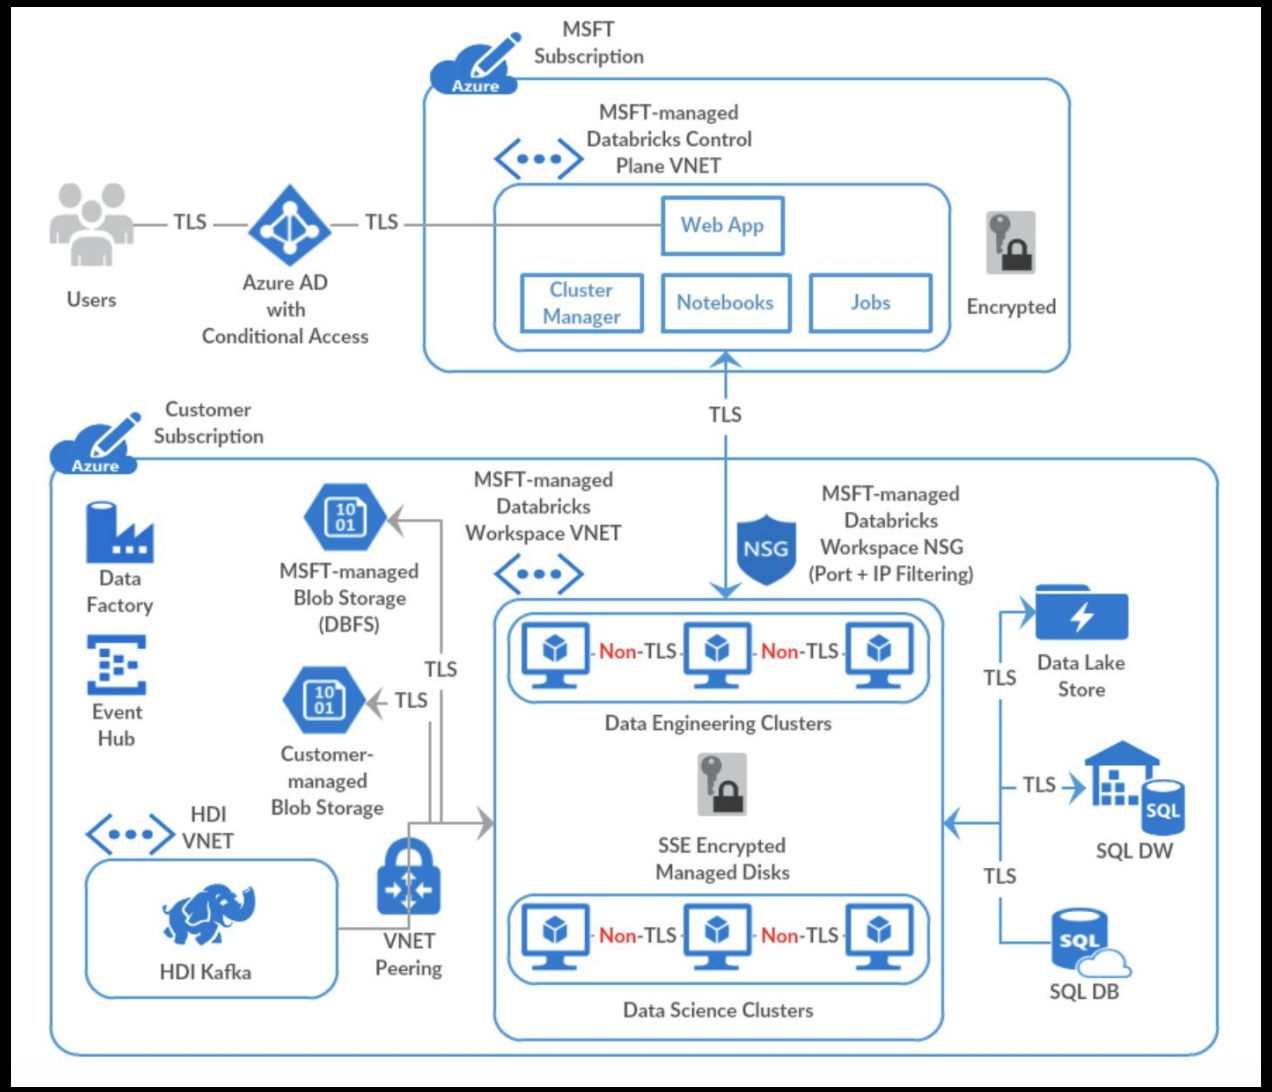

databricks

Azure Databricks Security

Azure Databricks Security

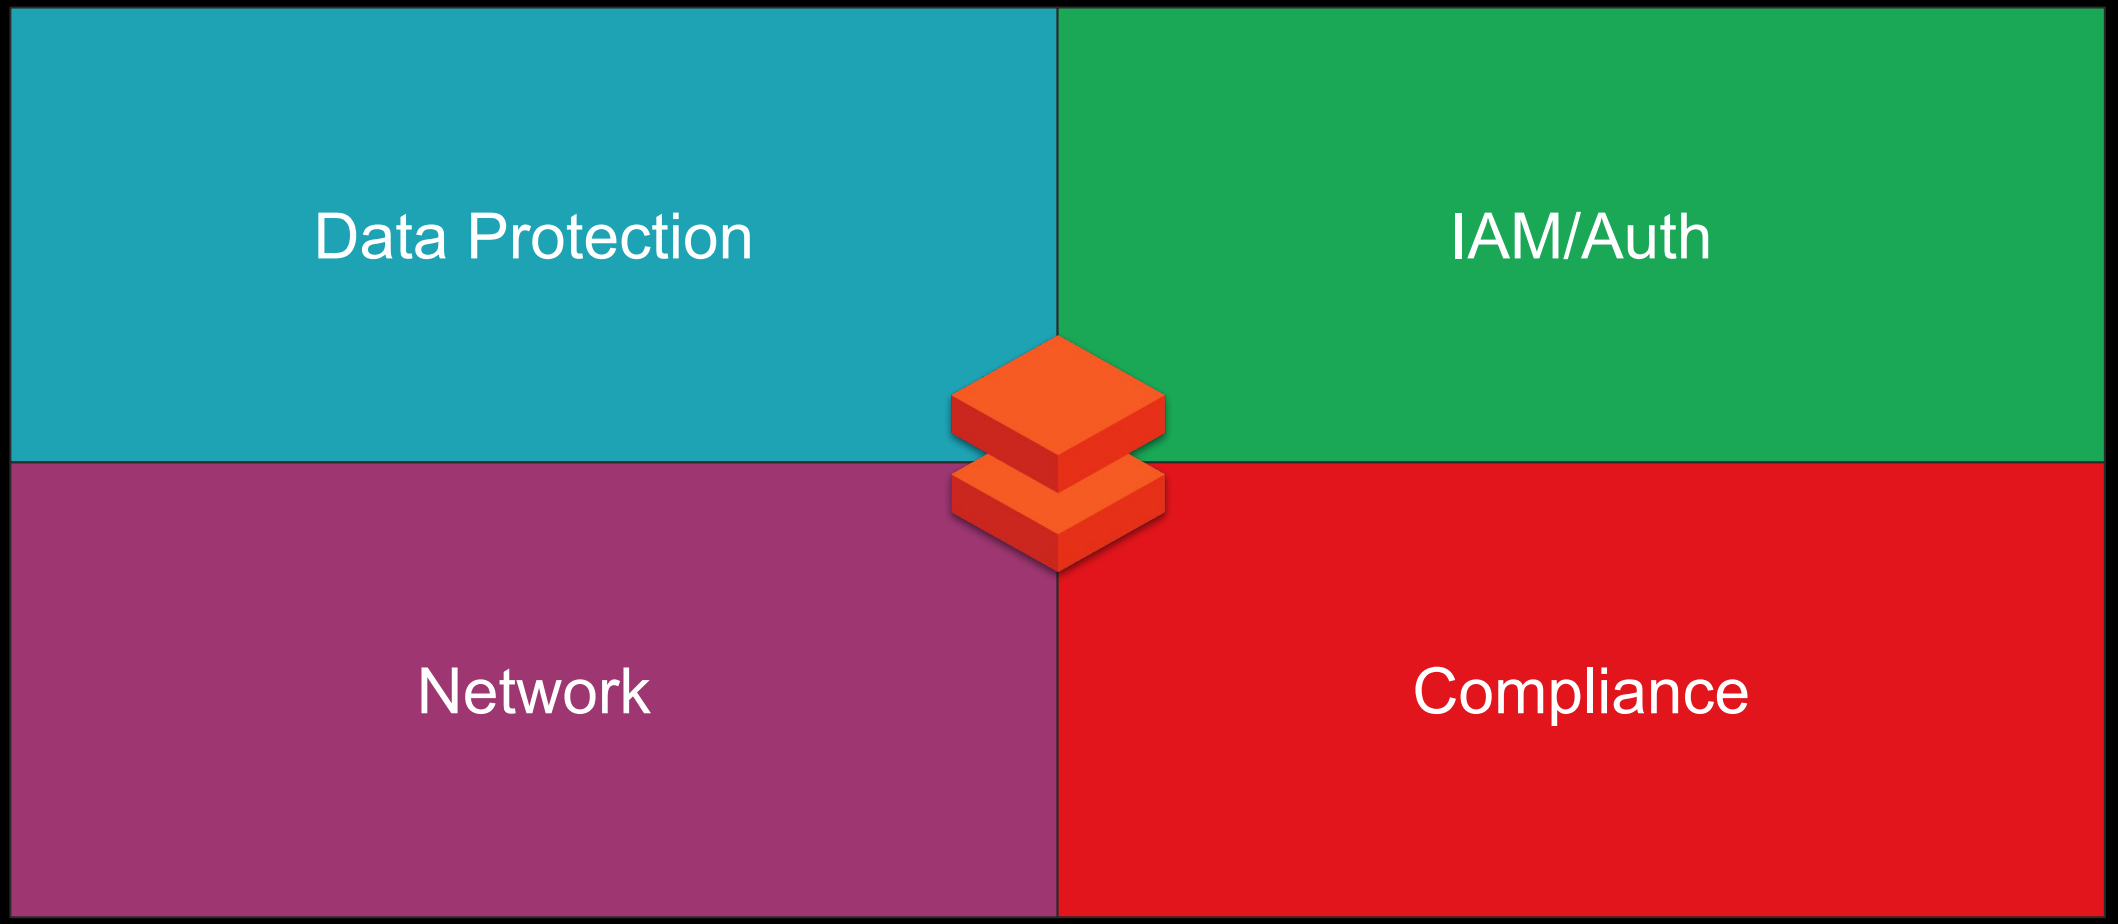

Azure Databricks Security | Data Protection

Data Protection

- Encryption-At-Rest – Service Managed Keys, User Managed Keys
- Encryption-in-flight (Transport Layer Security TLS)
- File/Folder Level ACLs for AAD Users, Groups, Service Principals
- ACLs for Clusters, Folders, Notebooks, Tables, Jobs
- Secrets with Azure Key Vault

* Preview feature

databricks

Data Protection | Encryption | At-Rest

- Azure Databricks has separation of compute and storage

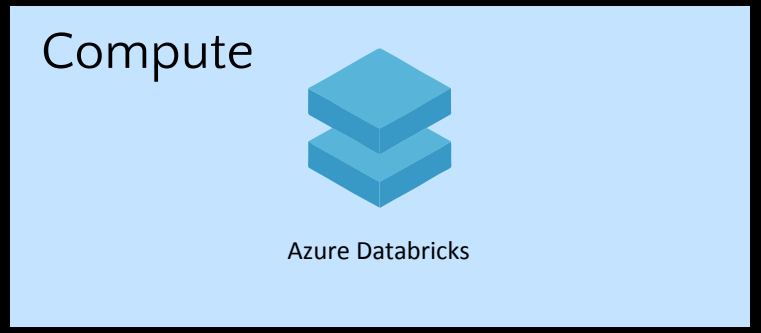

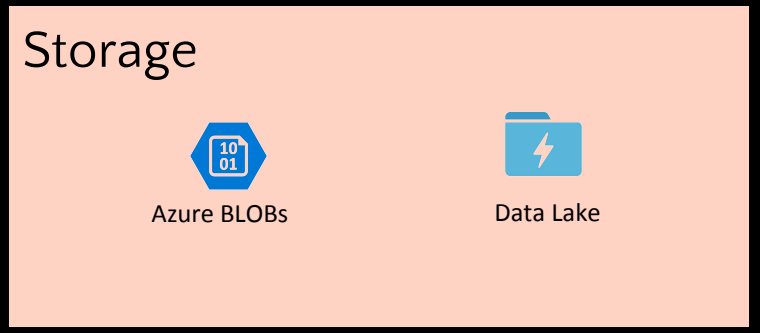

- Storage Services such as Azure Blob Store, Azure Data Lake Storage Provide
  - Encryption of Data
  - Customer Managed Keys
  - File/Folder Level ACLs (Azure Data Lake Storage)

Data Protection | Encryption | In-Transit

All the traffic from the Control Plane to the Clusters in the customer subscription is always encrypted with TLS.

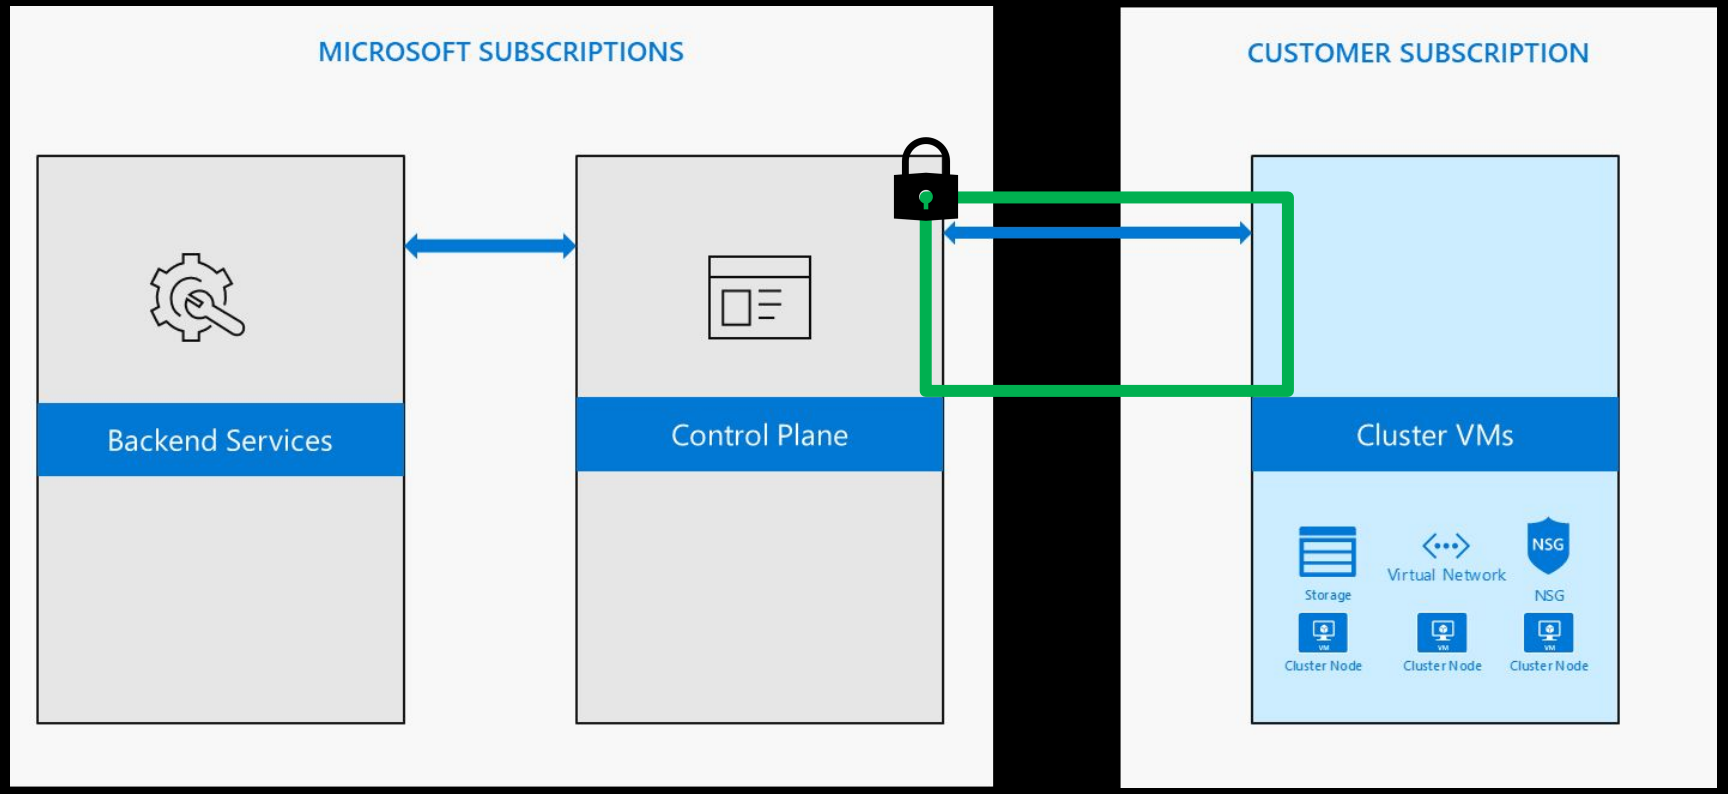

Data Protection | Access Control | ADLS Passthru

- Authenticate automatically to Azure Data Lake Storage (ADLS) from Azure Databricks clusters using the same Azure Active Directory (Azure AD) identity that one uses to log into Azure Databricks.
- Commands running on a configured cluster will be able to read and write data in Azure Data Lake Storage without requiring one to configure service principal credentials.

Data Protection | Access Control | Folders

Ability	No Permissions	Read	Run	Edit	Manage
View items		x	x	x	x
Create, clone, import, export items		x	x	x	x
Run commands on notebooks			x	x	x
Attach/detach notebooks			x	x	x
Delete items				x	x
Move/rename items				x	x
Change permissions					x

Data Protection | Access Control | Notebooks

Ability	No Permissions	Read	Run	Edit	Manage
View cells		x	x	x	x
Comment		x	x	x	x
Run commands			x	x	x
Attach/detach notebooks			x	x	x
Edit cells				x	x
Change permissions					x

Data Protection | Access Control | Clusters

Ability	No Permissions	Can Attach To	Can Restart	Can Manage
Attach notebook to cluster		x	x	x
View Spark UI		x	x	x
View cluster metrics		x	x	x
Terminate cluster			x	x
Start cluster			x	x
Restart cluster			x	x
Edit cluster				x
Attach library to cluster				x
Resize cluster				x
Modify permissions				x

Data Protection | Access Control | Jobs

Ability	No Permissions	Can View	Can Manage Run	Is Owner	Can Manage (admin)
View job details and settings	x	x	x	x	x
View results, Spark UI, logs of a job run		x	x	x	x
Run now			x	x	x
Cancel run			x	x	x
Edit job settings				x	x
Modify permissions				x	x

Data Protection | Access Control | Tables

Objects

CATALOG | DATABASE | TABLE | VIEW | FUNCTION | ANONYMOUS FUNCTION | ANY FILE

Privileges

SELECT – read access to an object
CREATE – ability to create an object (eg. Table in a Database)
MODIFY – ability to add/delete/modify data in an Object
READ_METADATA – ability to read Metadata about an object
ALL_PRIVELEGES – all of the above

Data Protection | Secrets

- Using our Secrets APIs, Secrets can be securely stored including in a Azure Key Vault or Databricks backend
- Authorized users can consume the secrets to access services

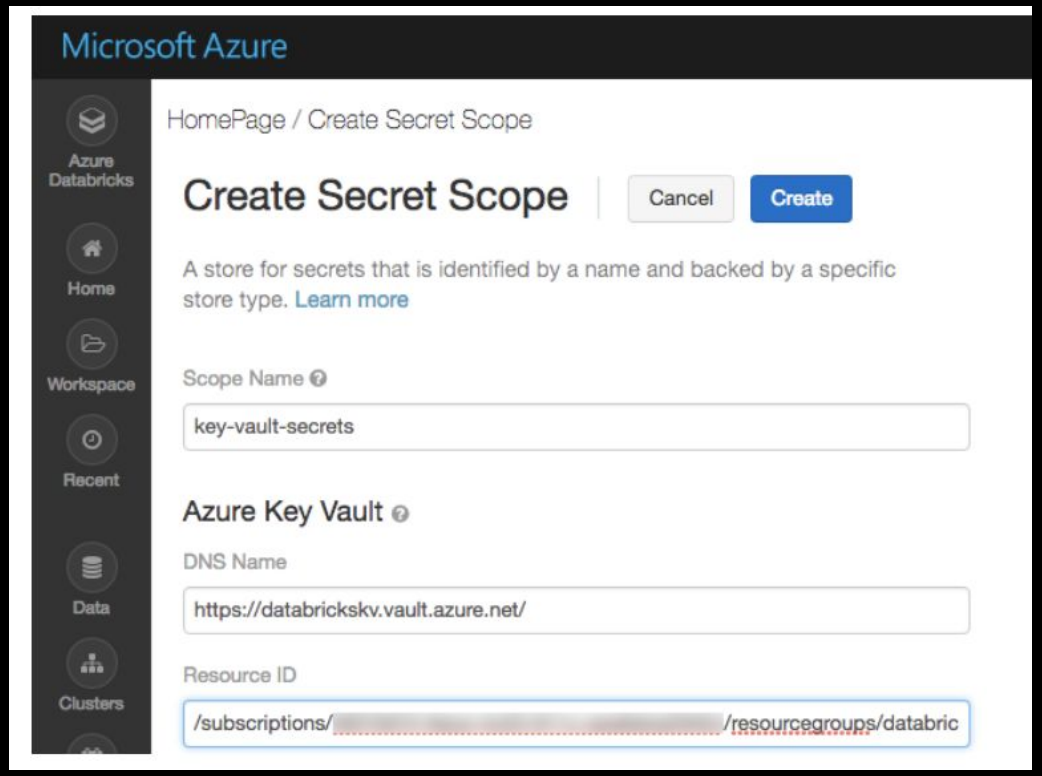

Azure Databricks Security | IAM/Auth

IAM/Auth

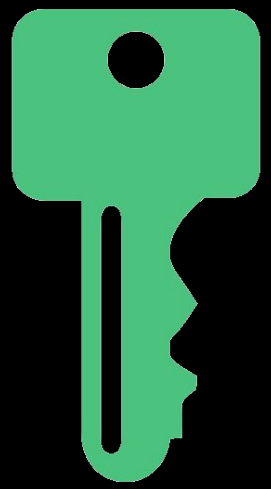

- Azure Active Directory (AAD) Authentication (w/ MFA)
- AAD Groups (using SCIM)
- AAD Conditional Access
- AAD Access Tokens

* Preview feature

IAM/Auth | First-party AAD Integration

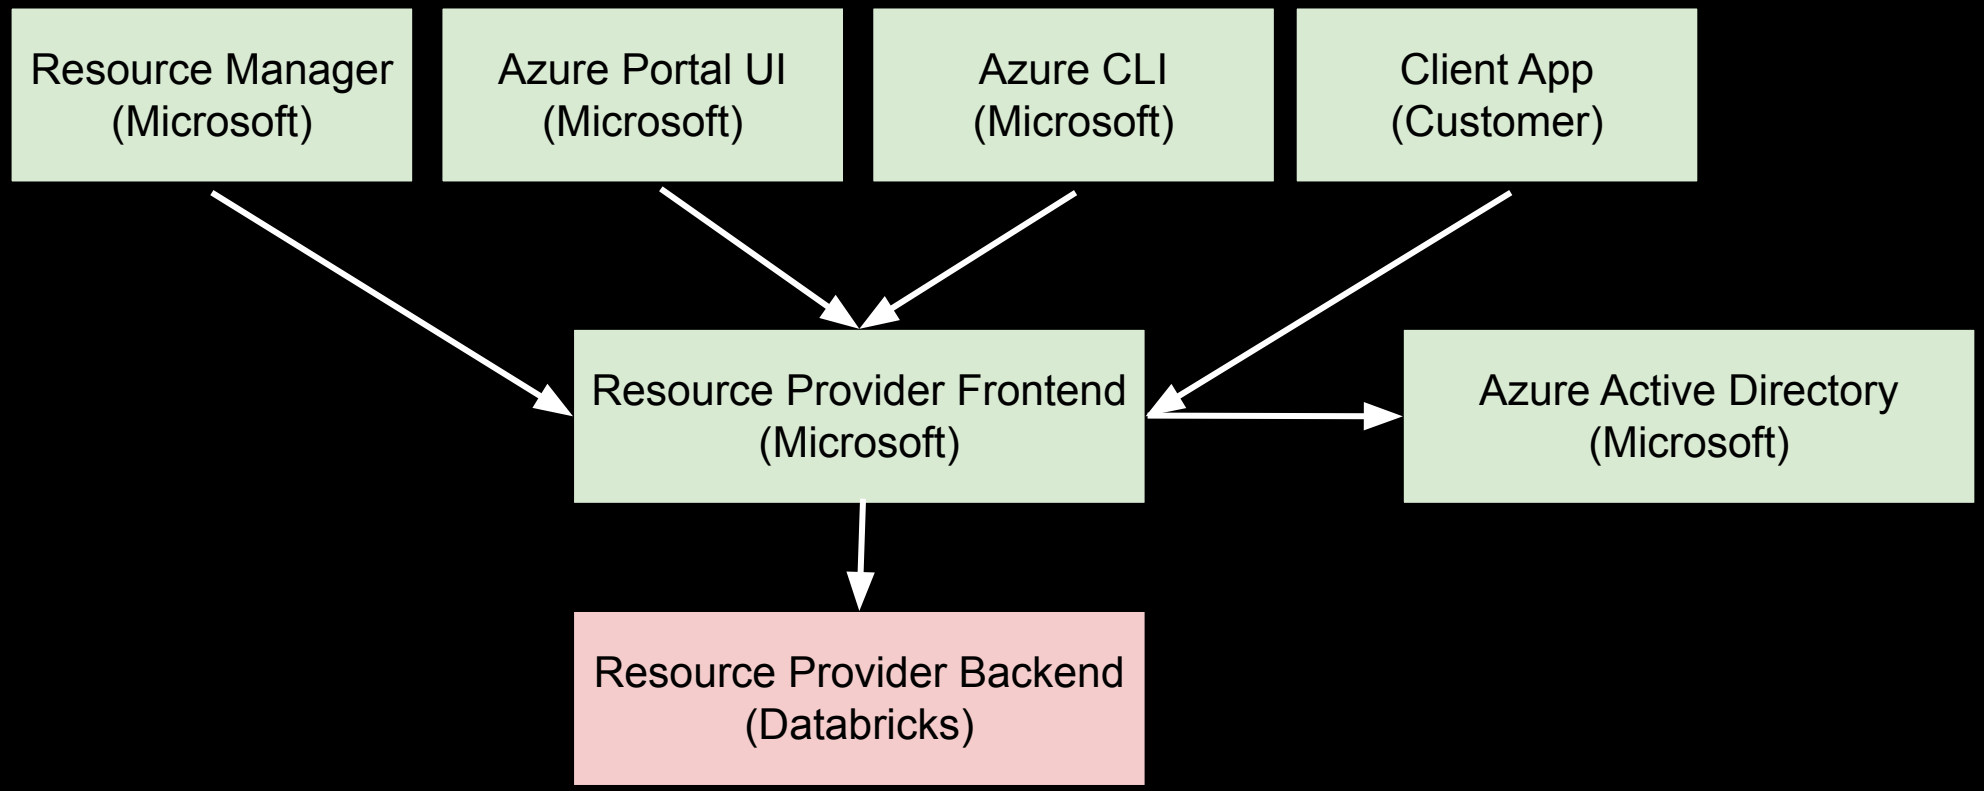

IAM/Auth | SCIM Integration

Azure Databricks supports SCIM, or System for Cross-domain Identity Management, an open standard that allows you to automate user provisioning. SCIM lets you use Azure Active Directory to create users in Azure Databricks and give them the proper level of access, as well as remove access for users (deprovision them) when they leave the organization or no longer need access to Azure Databricks.

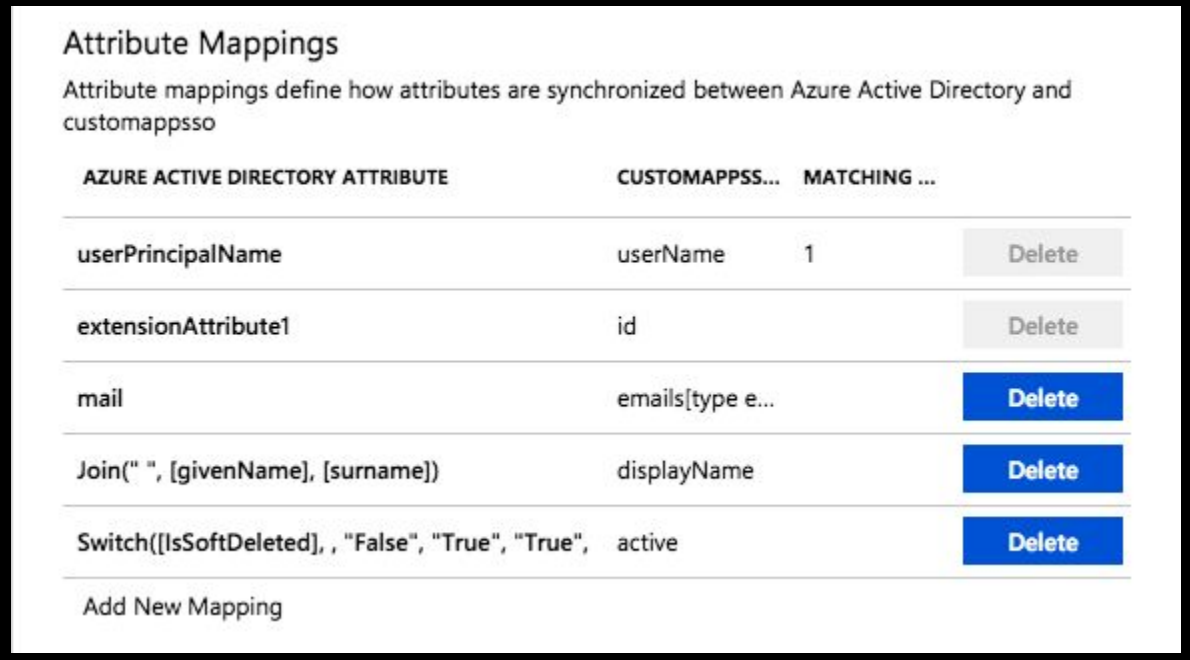

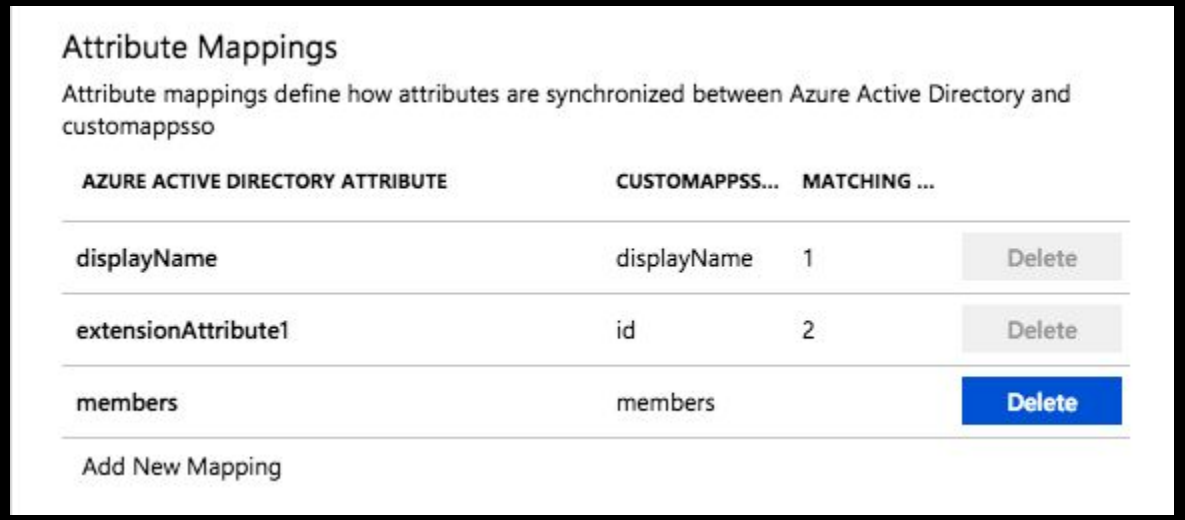

databricks

IAM/Auth | Conditional Access

Azure Databricks supports Azure Active Directory conditional access, which allows administrators to control where and when users are permitted to sign in to Azure Databricks. For example, conditional access policies can restrict sign-in to your corporate network or can require multi-factor authentication.

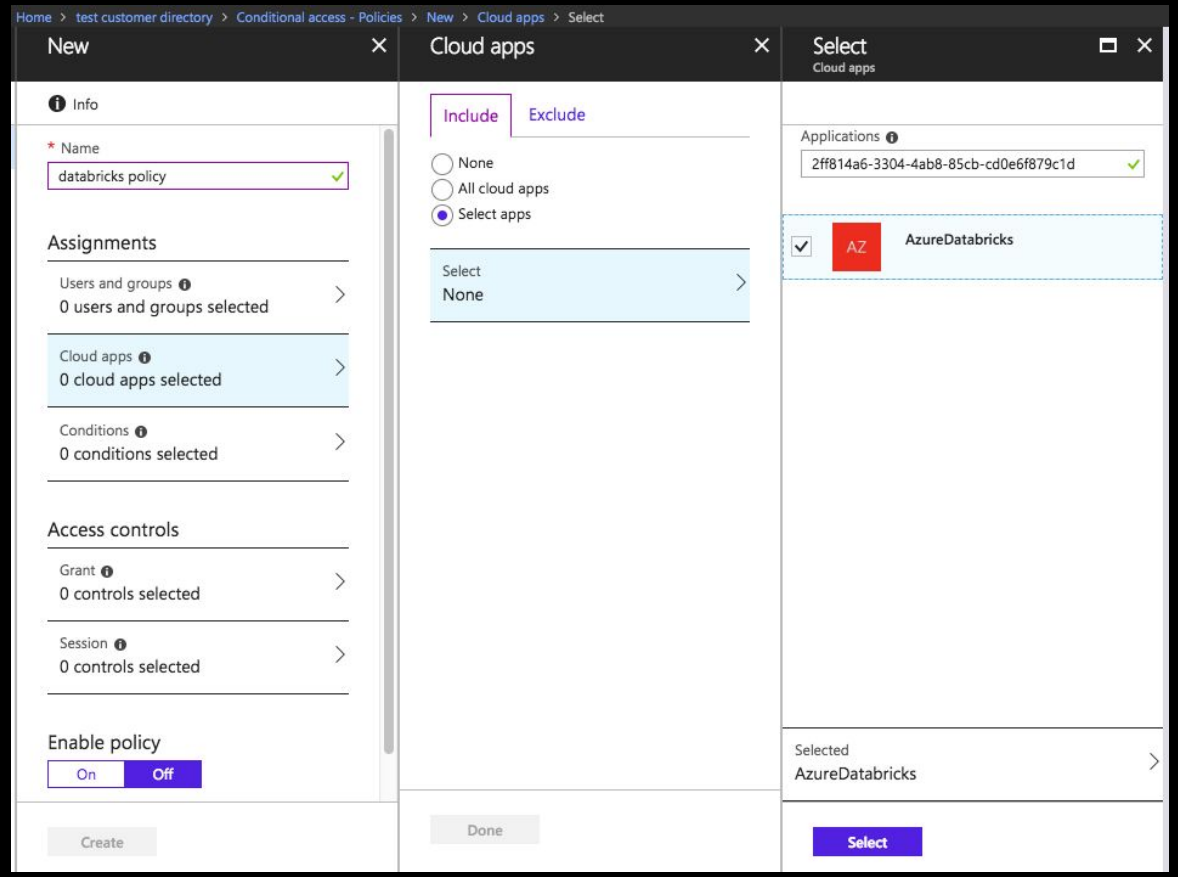

databricks

IAM/Auth | AAD Token Support

You could use AAD tokens to automate provisioning of Azure Databricks workspaces and access the Databricks REST API

****Private Preview Only****

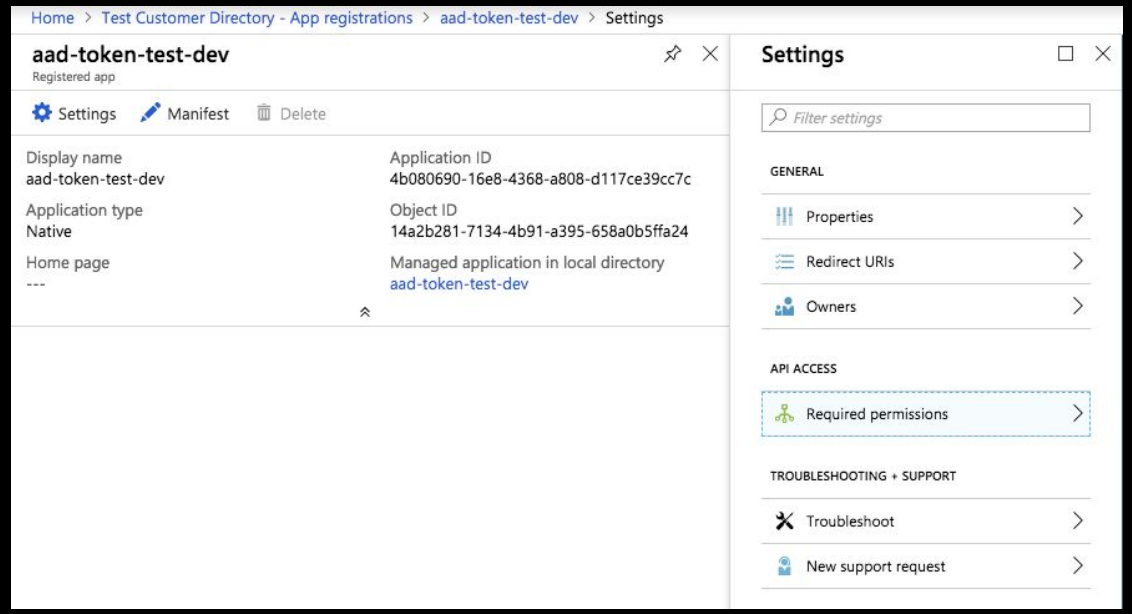

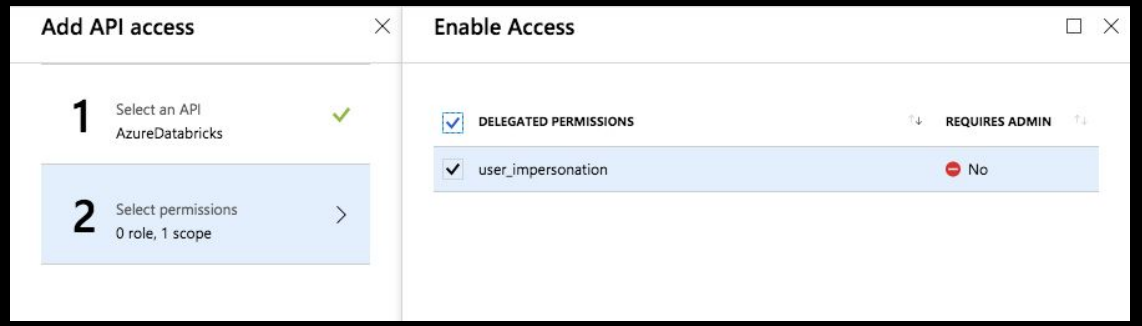

databricks

Azure Databricks Security | Network Security

Network Security

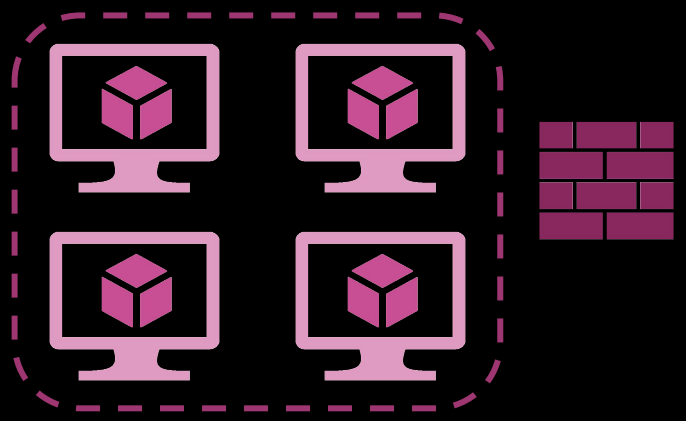

- Managed VNets
- VNet Peering
- VNET Injection/BYO VNET
  - On-Premises Data Access
  - Single-IP SNAT and Firewall-based filtering via custom routing
  - Service Endpoint

* Preview feature

Network Security | VNET Peering

Virtual network (VNet) peering allows the virtual network in which your Azure Databricks resource is running to peer with another Azure virtual network. Traffic between virtual machines in the peered virtual networks is routed through the Microsoft backbone infrastructure, much like traffic is routed between virtual machines in the same virtual network, through private IP addresses only.

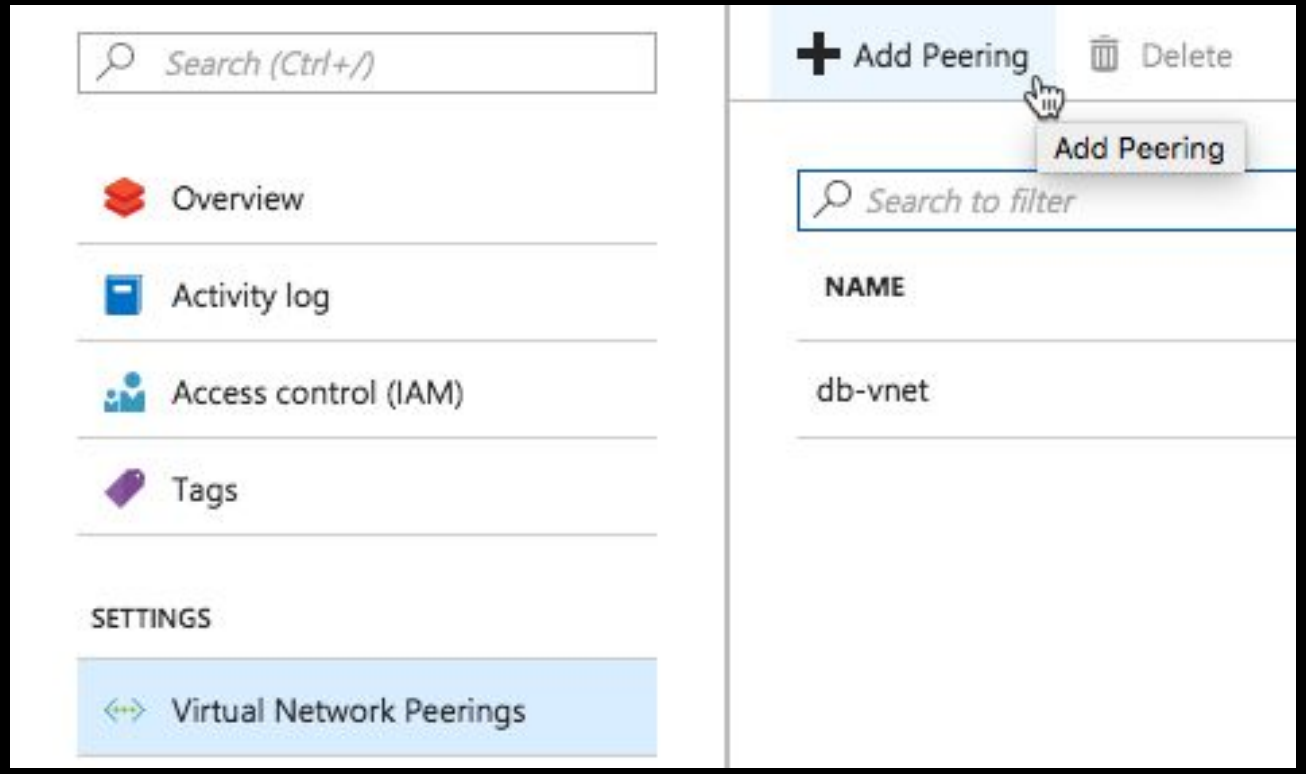

databricks

Network Security | VNET Injection / BYO VNET

If you're looking to do specific network customizations, you could deploy Azure Databricks data plane resources in your own VNET

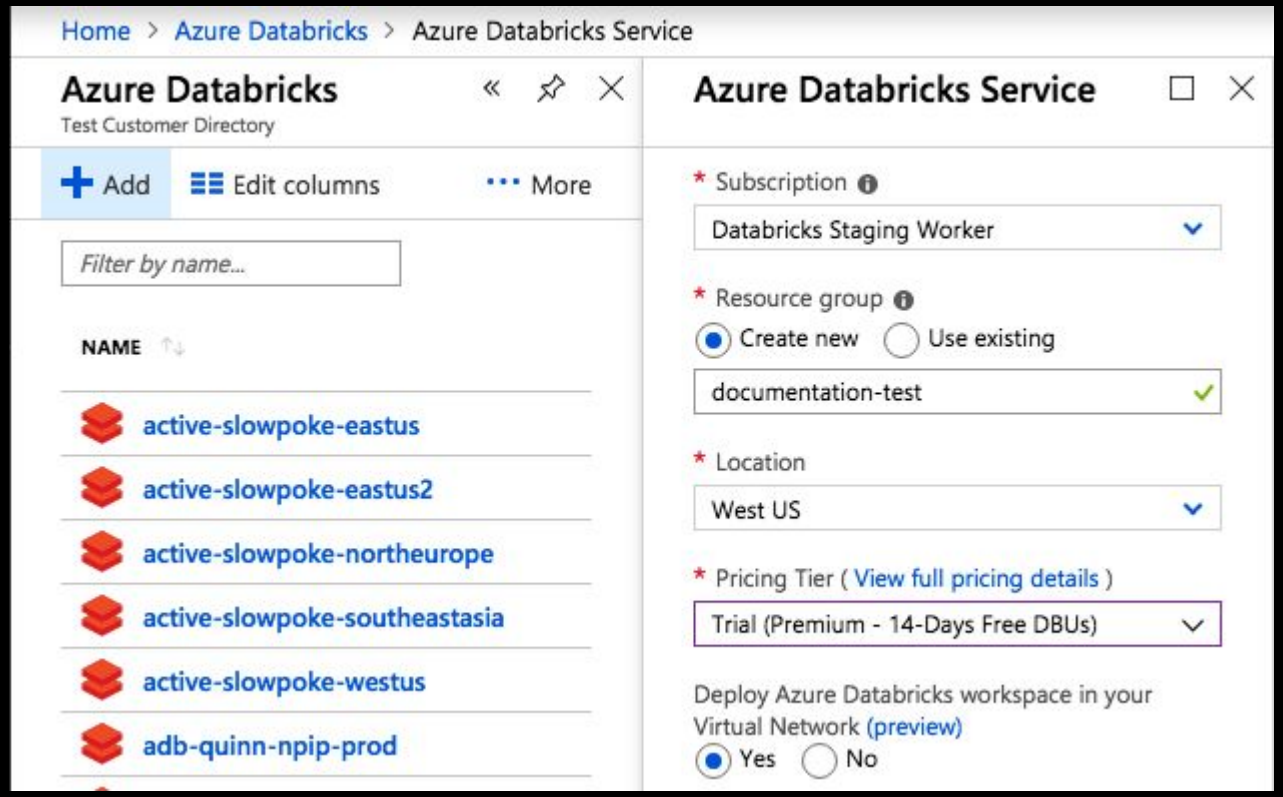

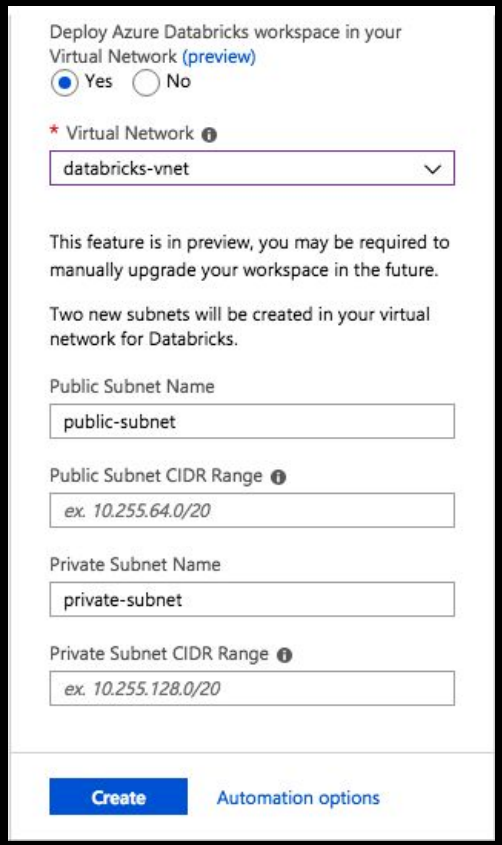

databricks

Azure Databricks Platform Architecture

Deployment with VNET Injection and inter-node TLS (one could be used without the other)

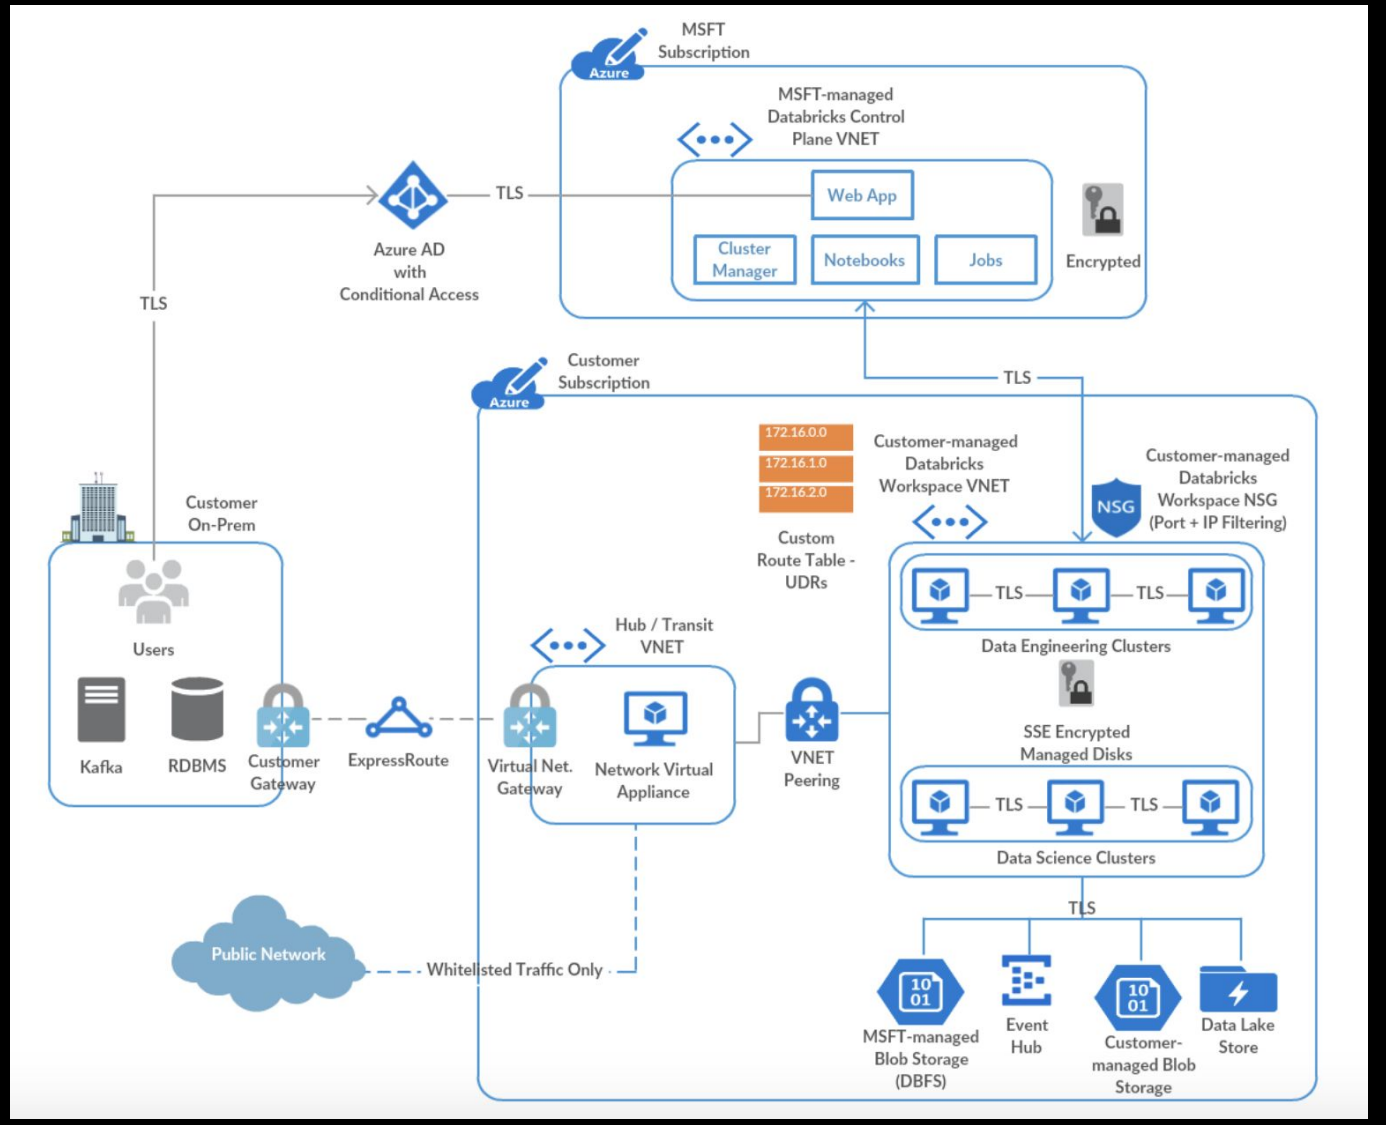

databricks

Private, secure, encrypted with compute in your network

Plug into your identity system for seamless control

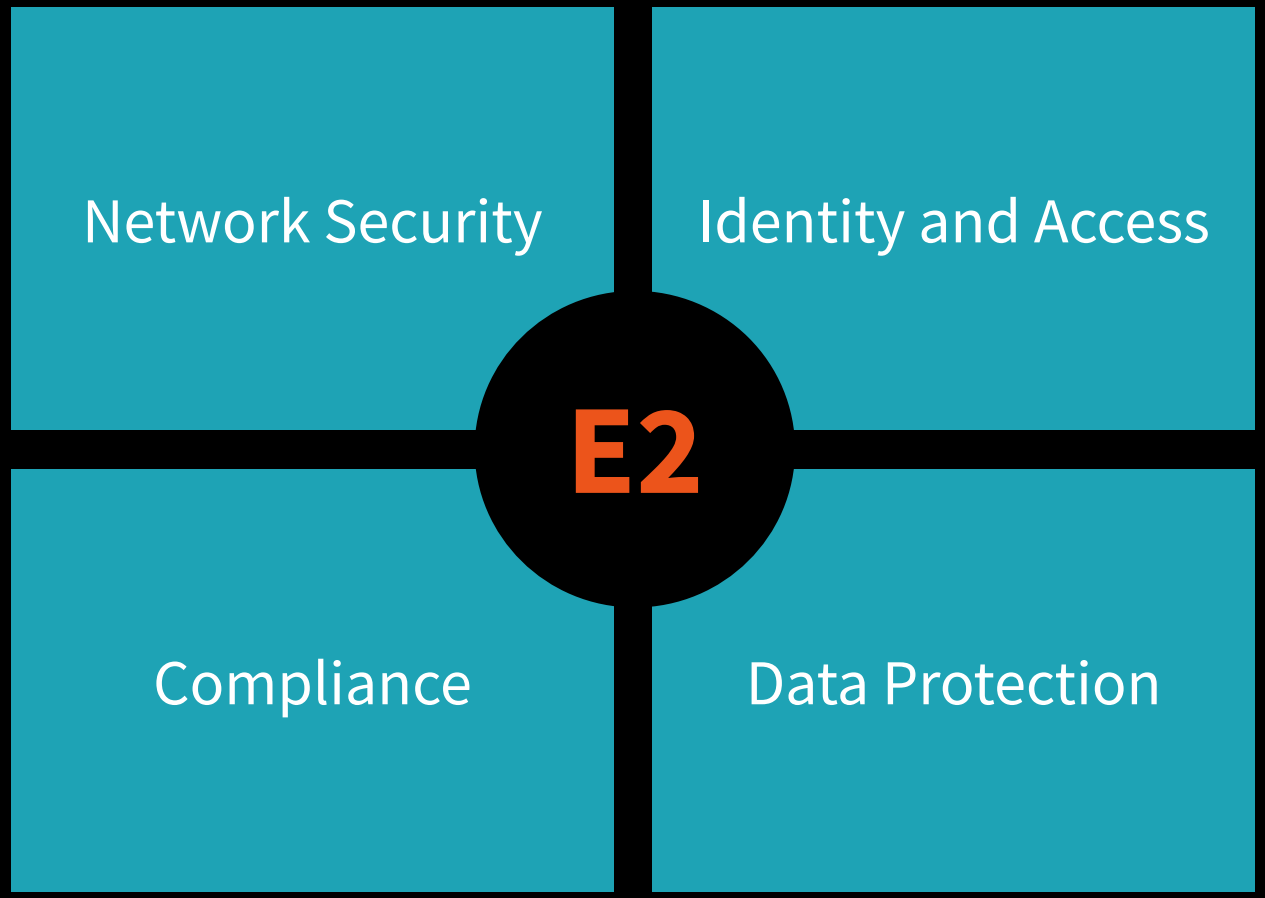

Battle tested for compliance and other sensitive data policies

Extend identity all the way down to the data natively

Secure network connectivity

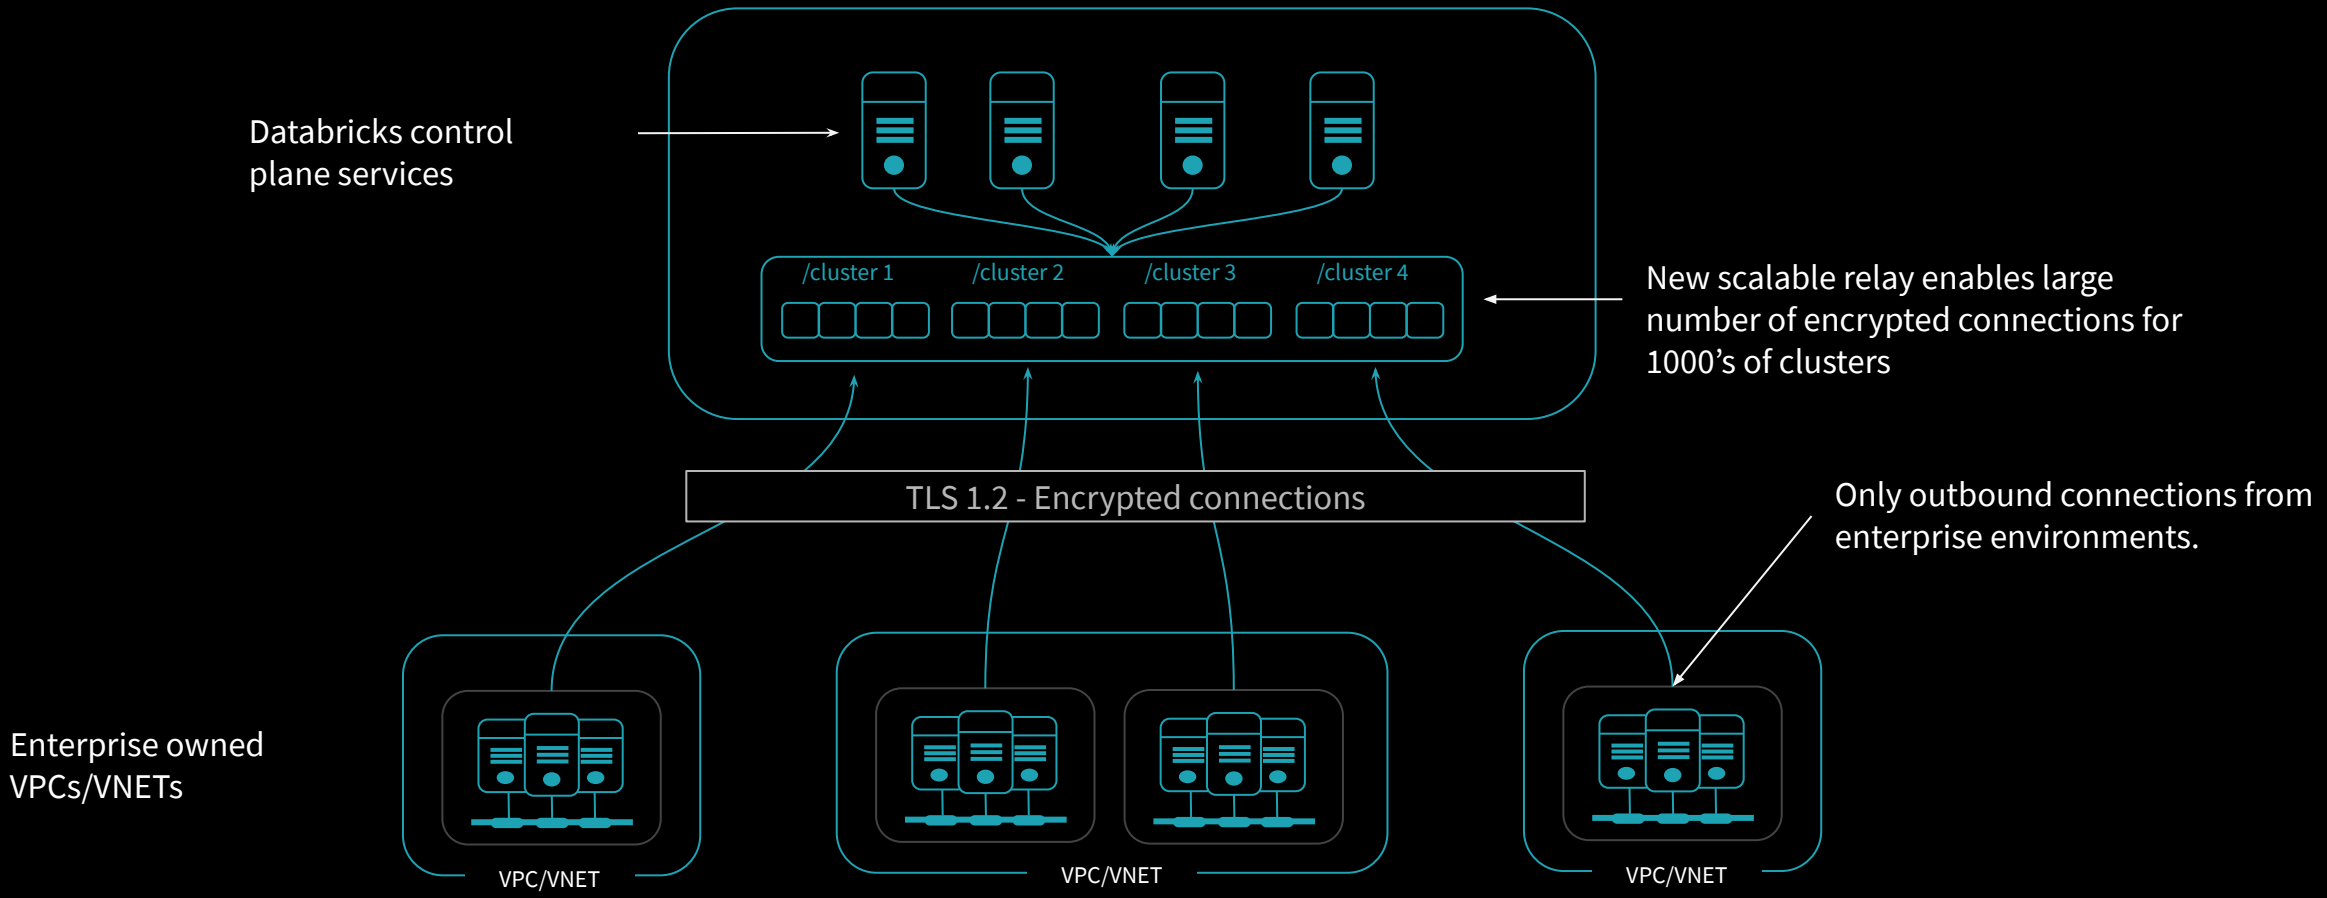

Private Link

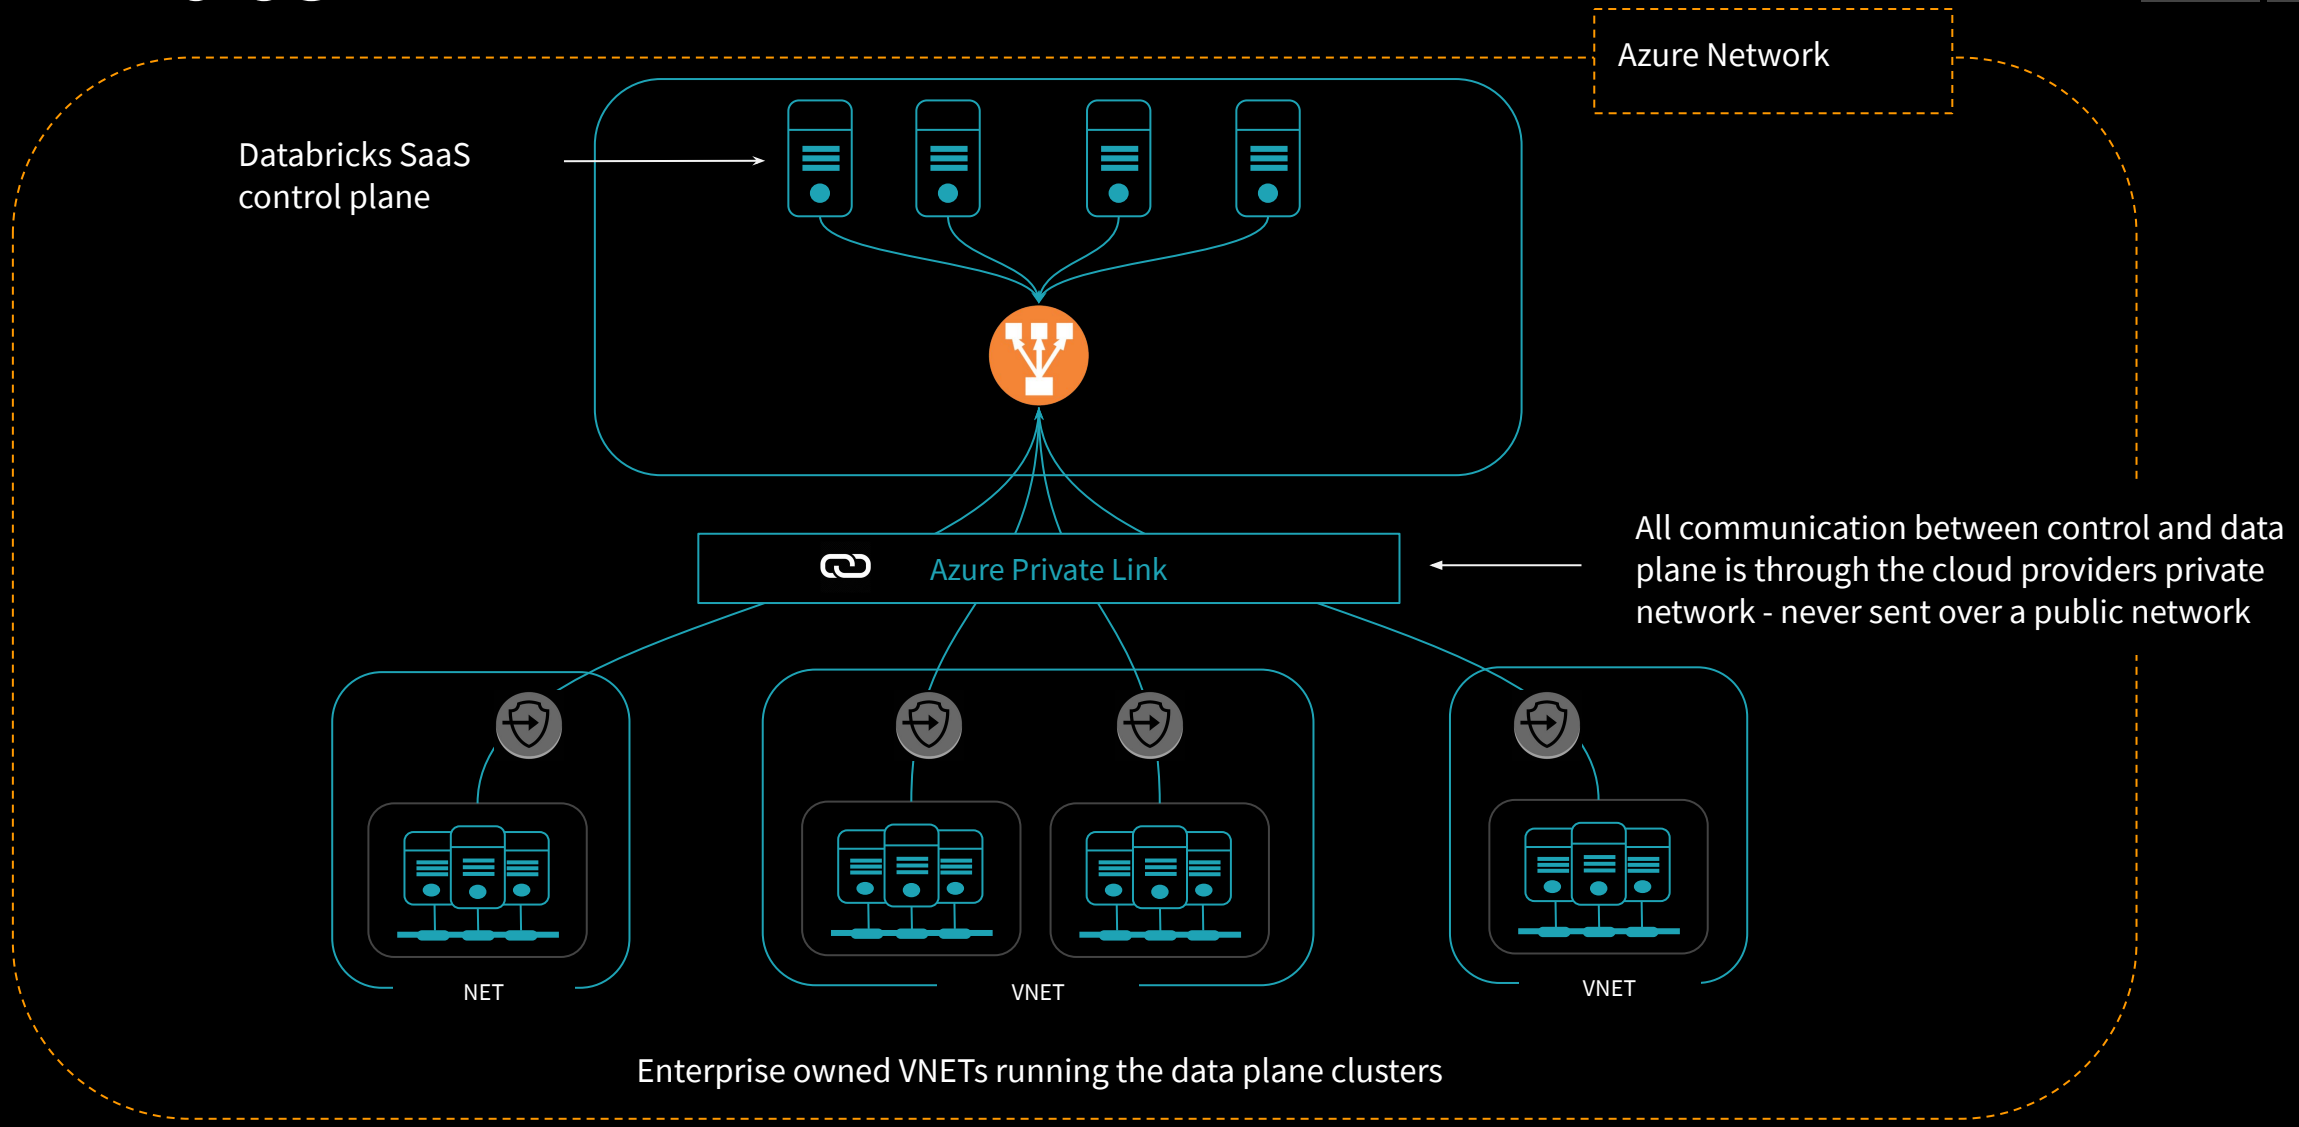

IP Access Lists

Control the networks that can access databricks

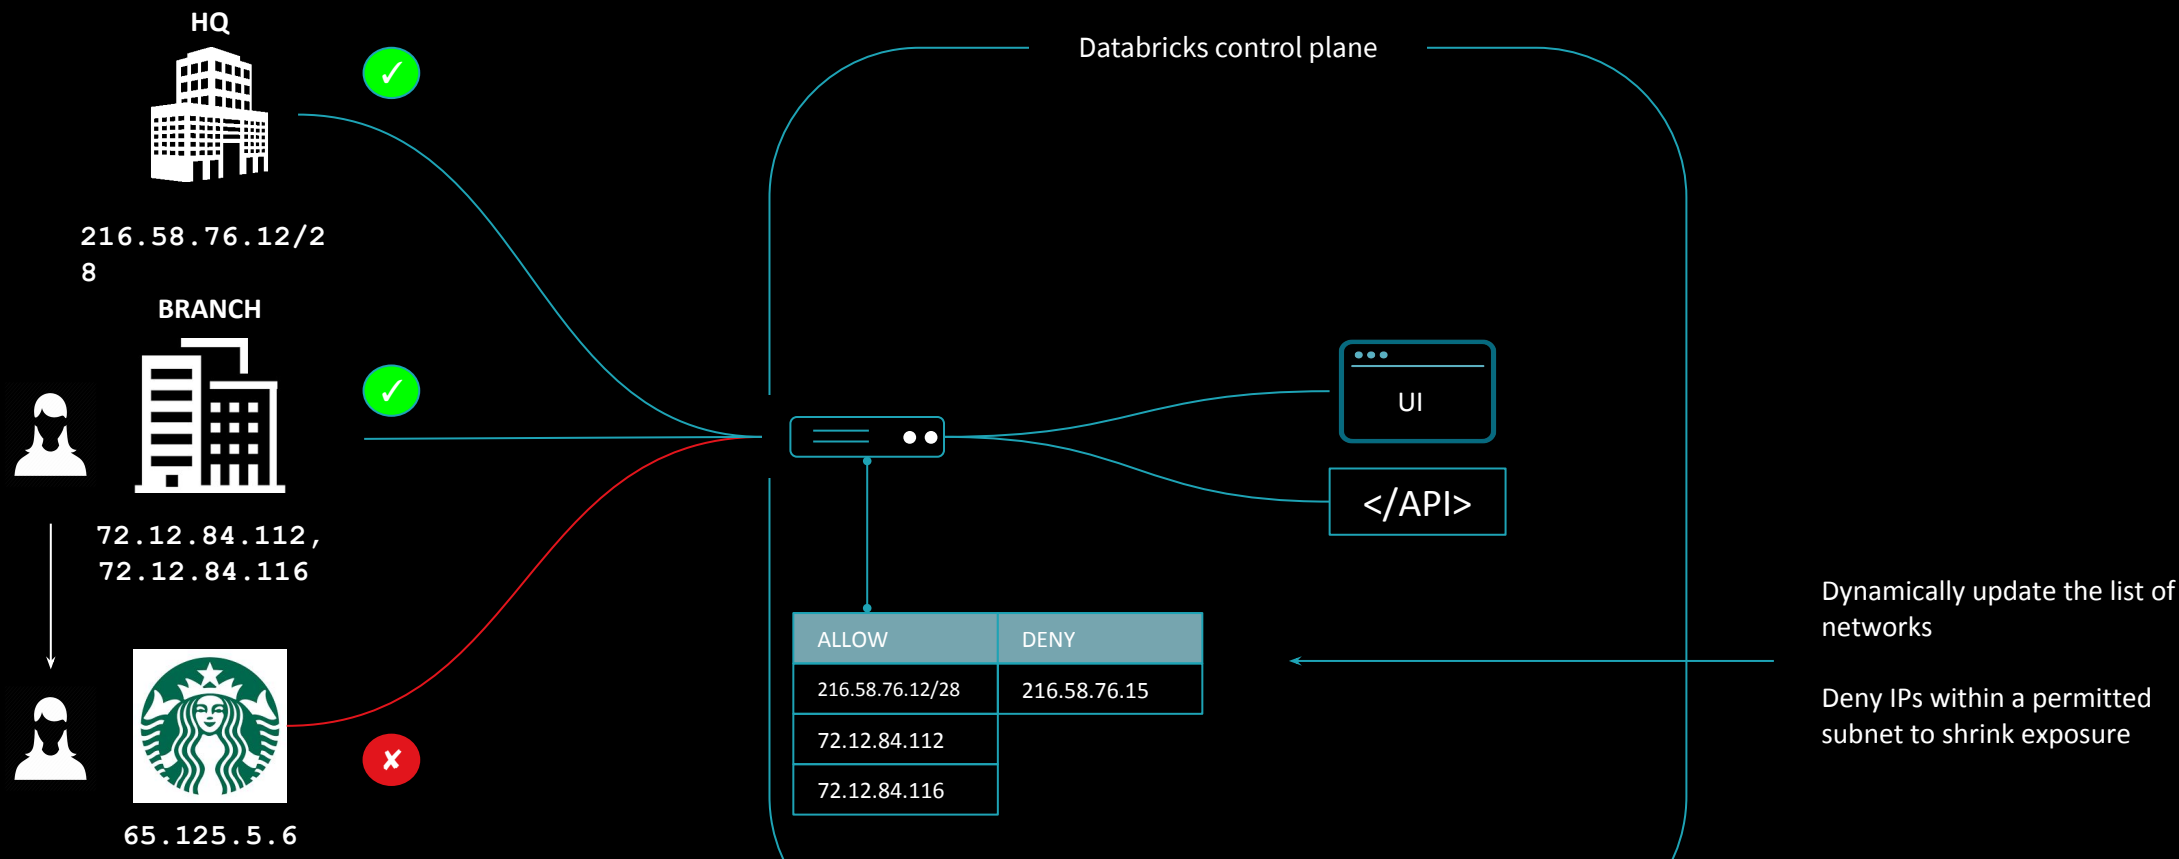

Bring your own VNET

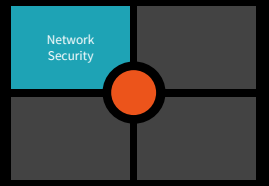

Today

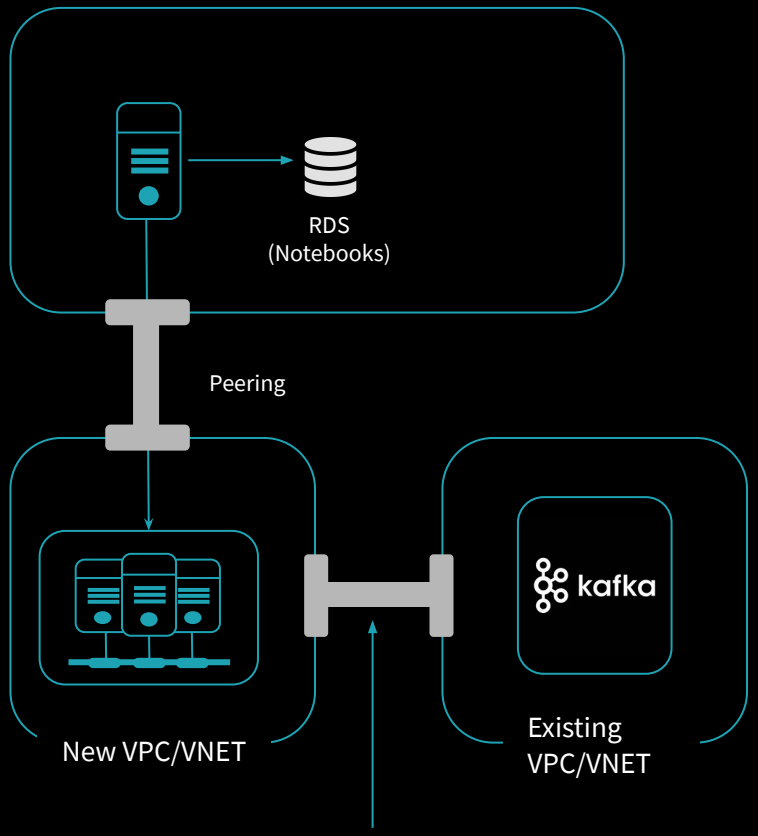

Connectivity to other applications through peering

Q4

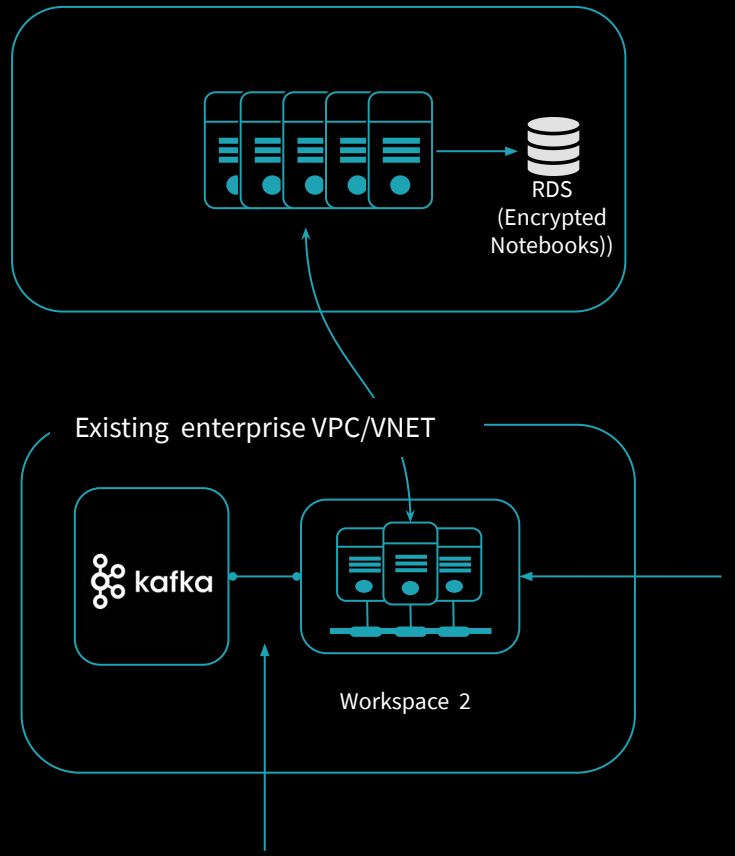

Databricks managed clusters in enterprise managed VPC/VNET

Local connectivity to other applications without peering

Data Centric Security

Example: End to end security for your data lake, defined in your metastore

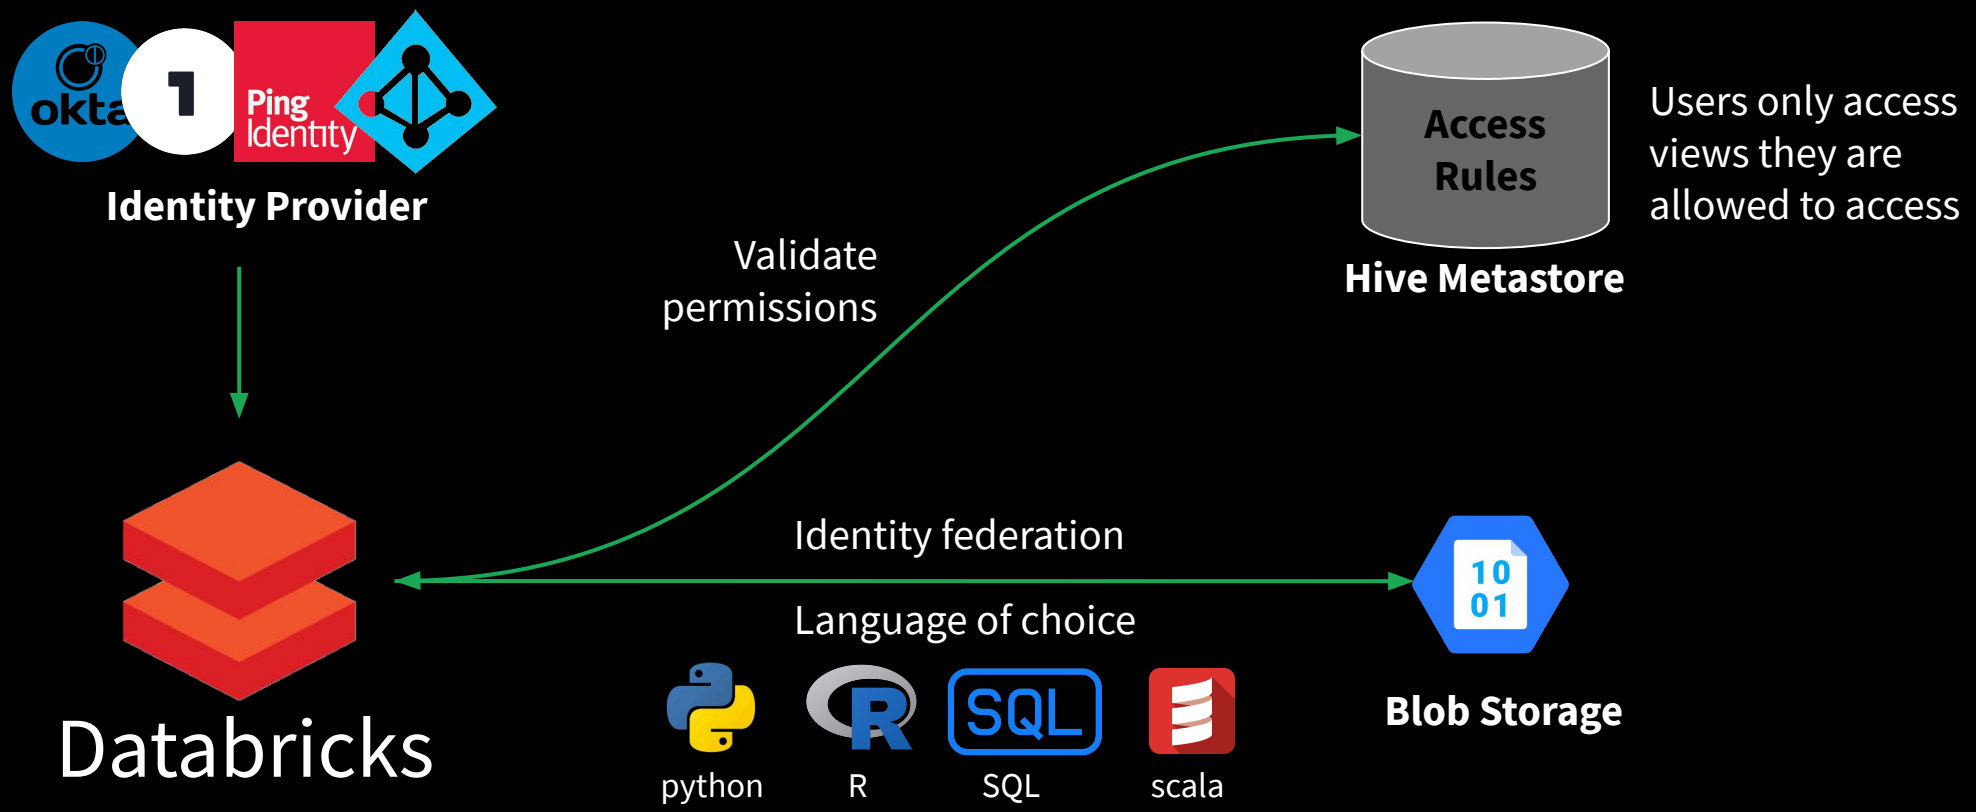

Azure Databricks

Data Centric Security

Example: End to end security for your data lake, defined in your metastore

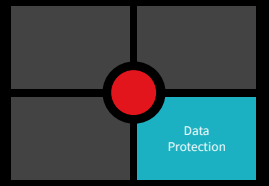

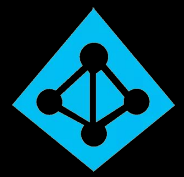

Azure Active Directory

Validate permissions

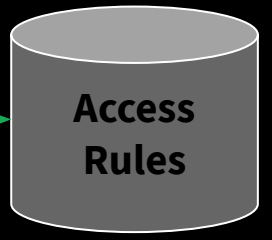

Hive Metastore

Users only access views they are allowed to access

Databricks

Identity federation

Language of choice

python

R

SQL

scala

Blob Storage

Azure Data Lake Gen2

Currently Preview

Encryption with enterprise root keys

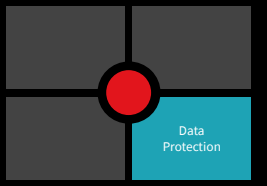

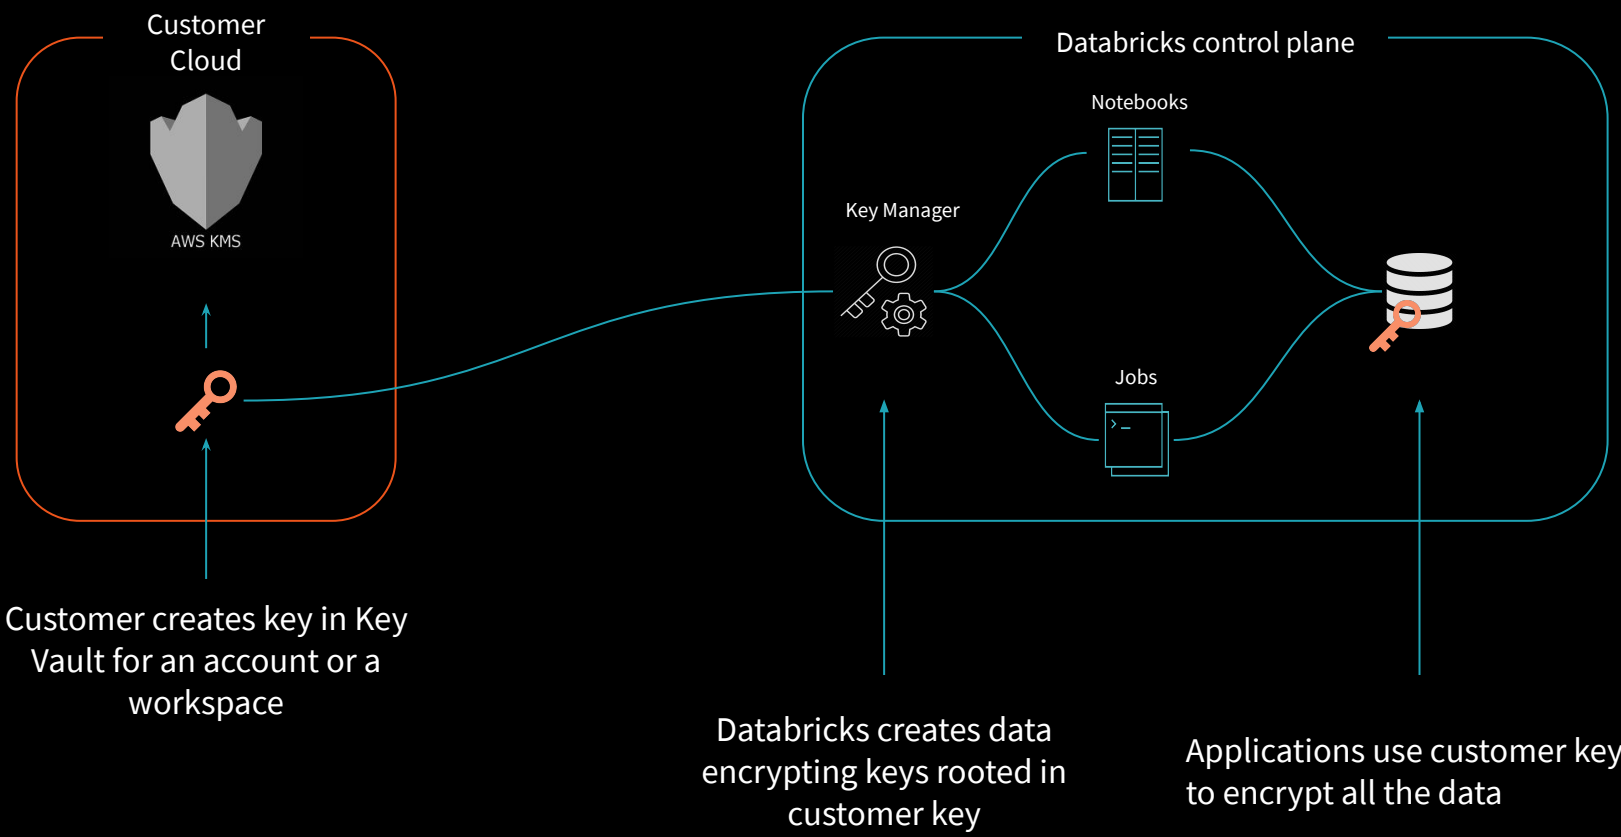

Customer creates key in Key Vault for an account or a workspace

Databricks creates data encrypting keys rooted in customer key

Applications use customer key to encrypt all the data

Get full control over keys used to encrypt data in the control plane. Revoking key revokes data access

Key hierarchy enables use of different keys for different notebooks

An audit log of key operations makes for easy reporting and audit requests

Currently Preview

End-to-end encrypted clusters

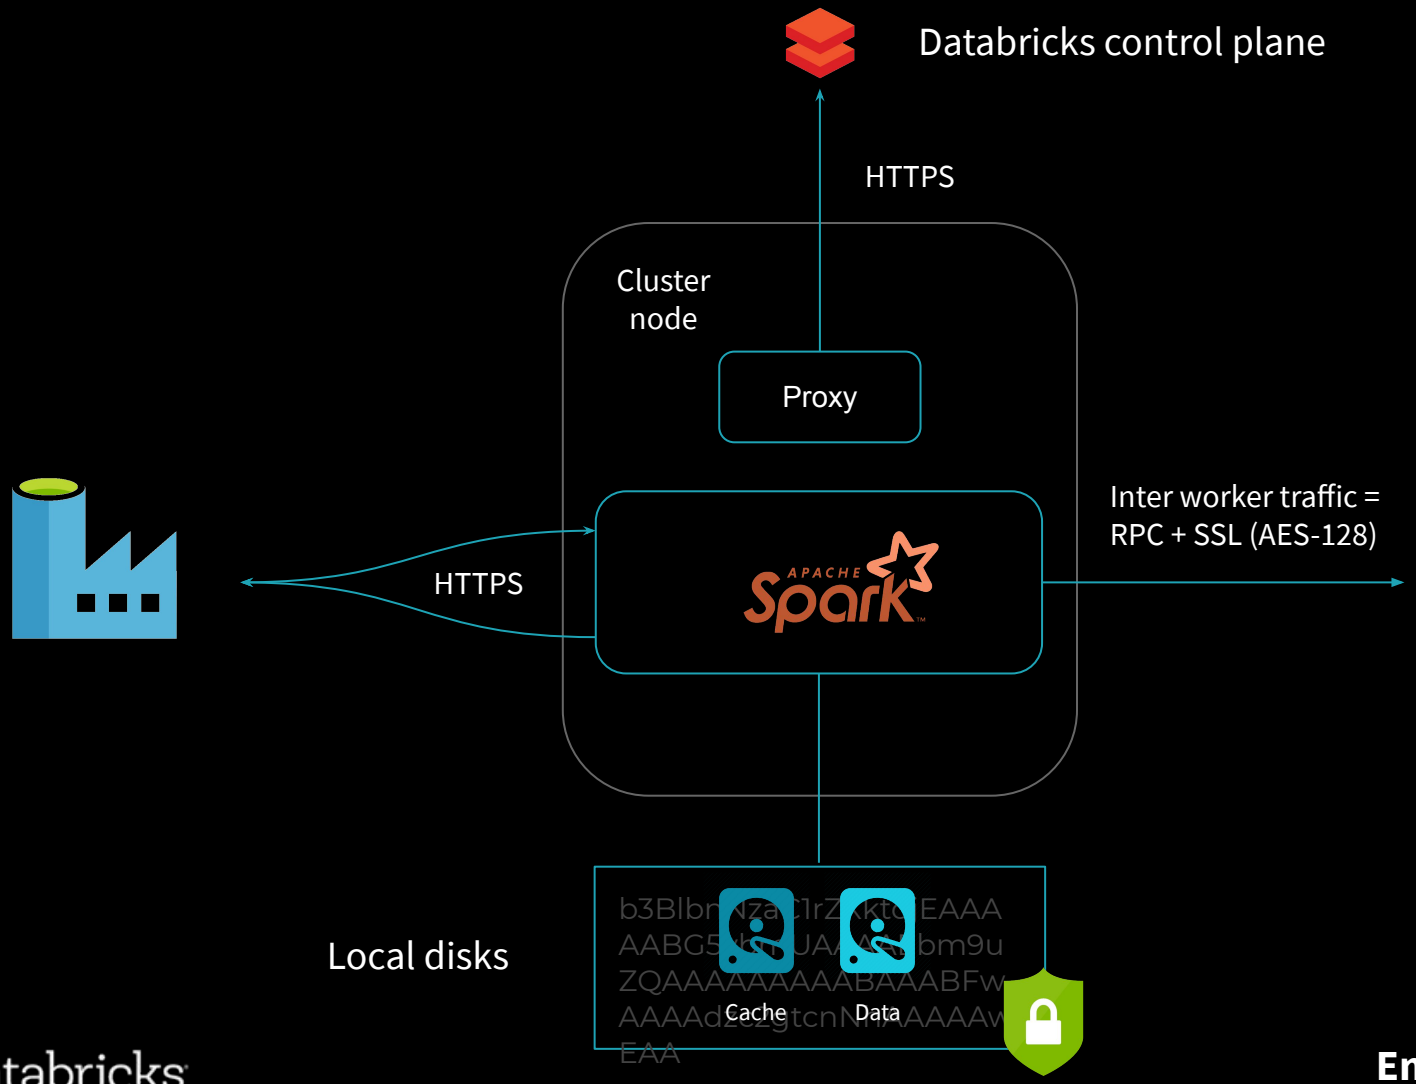

Enable encryption for a cluster - no key/cert management necessary

Ensure that data shared between cluster nodes is always encrypted

Guarantee that data is always encrypted no matter where its stored on the cluster - root, ephemeral or network attached disks

Security and Privacy

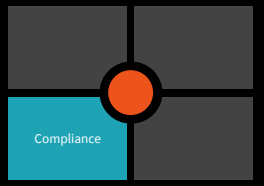

Security

1. ISO 27001
2. SOC 2 Type 2
3. HITRUST (end of 2019 on Azure)
4. GovCloud & FedRAMP (2020 on Azure)

Privacy

1. Privacy Shield
2. ISO 27018
3. GDPR compliant

Azure Databricks Security

Compliance

- Audit Logs
- ISO 27001
- ISO 27018
- HIPAA (Covered by MSFT BAA)
- SOC2, Type 2

* Preview feature

Compliance | Audit Logs

Databricks provides comprehensive end-to-end audit logs of activities performed by Databricks users, allowing your enterprise to monitor detailed Databricks usage patterns.

****Integration with Azure Monitor****

Services / Entities included are:

- Accounts
- Clusters
- DBFS
- Genie
- Jobs
- ACLs
- SSH
- Tables

Reference Architectures

Recommended End-to-End Architecture

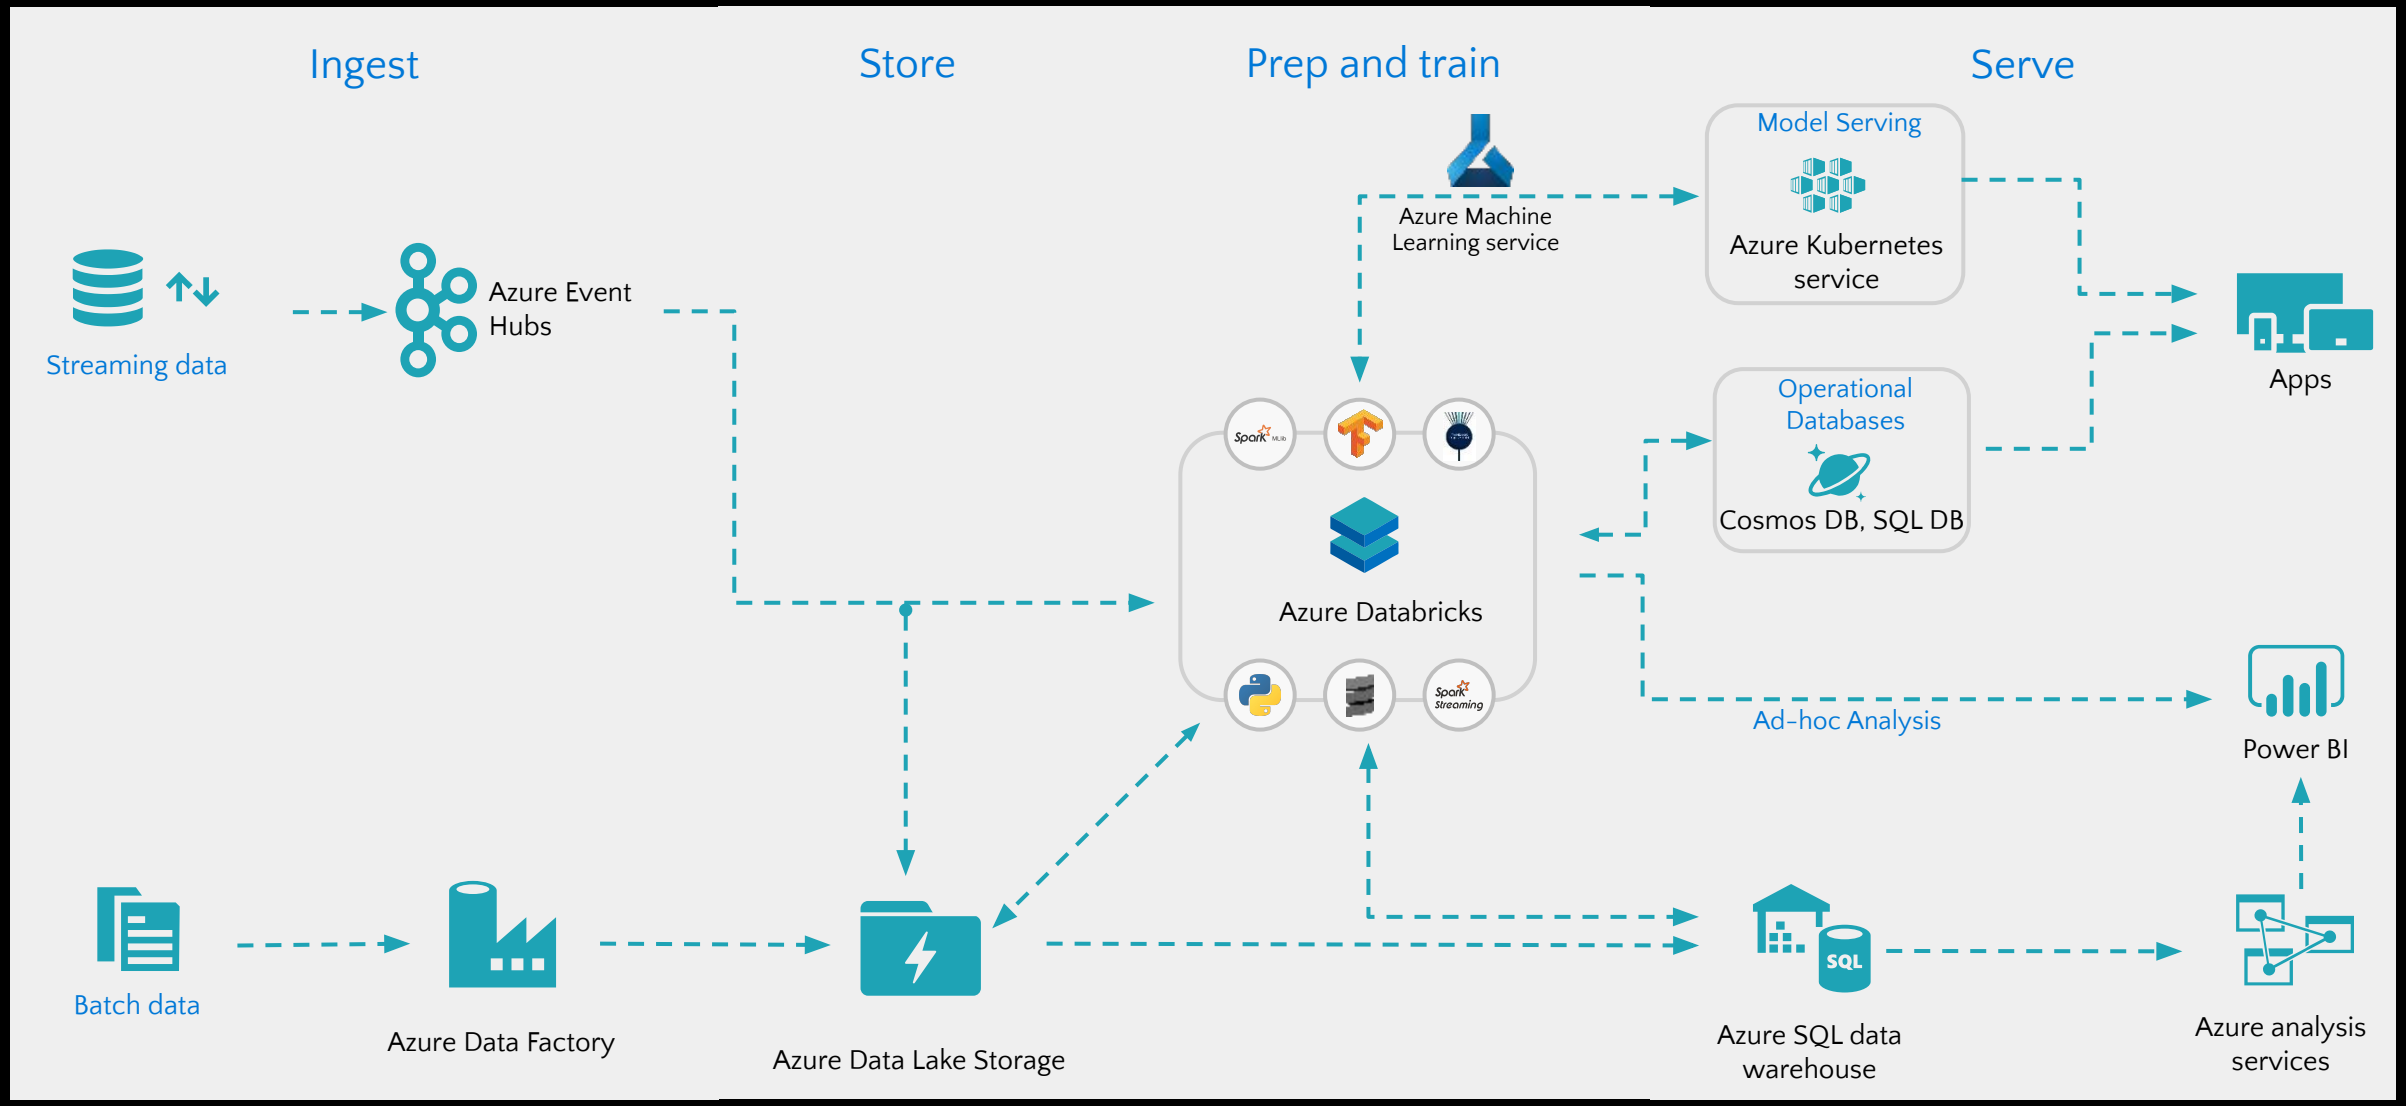

databricks

Recommended End-to-End Architecture - Batch ETL

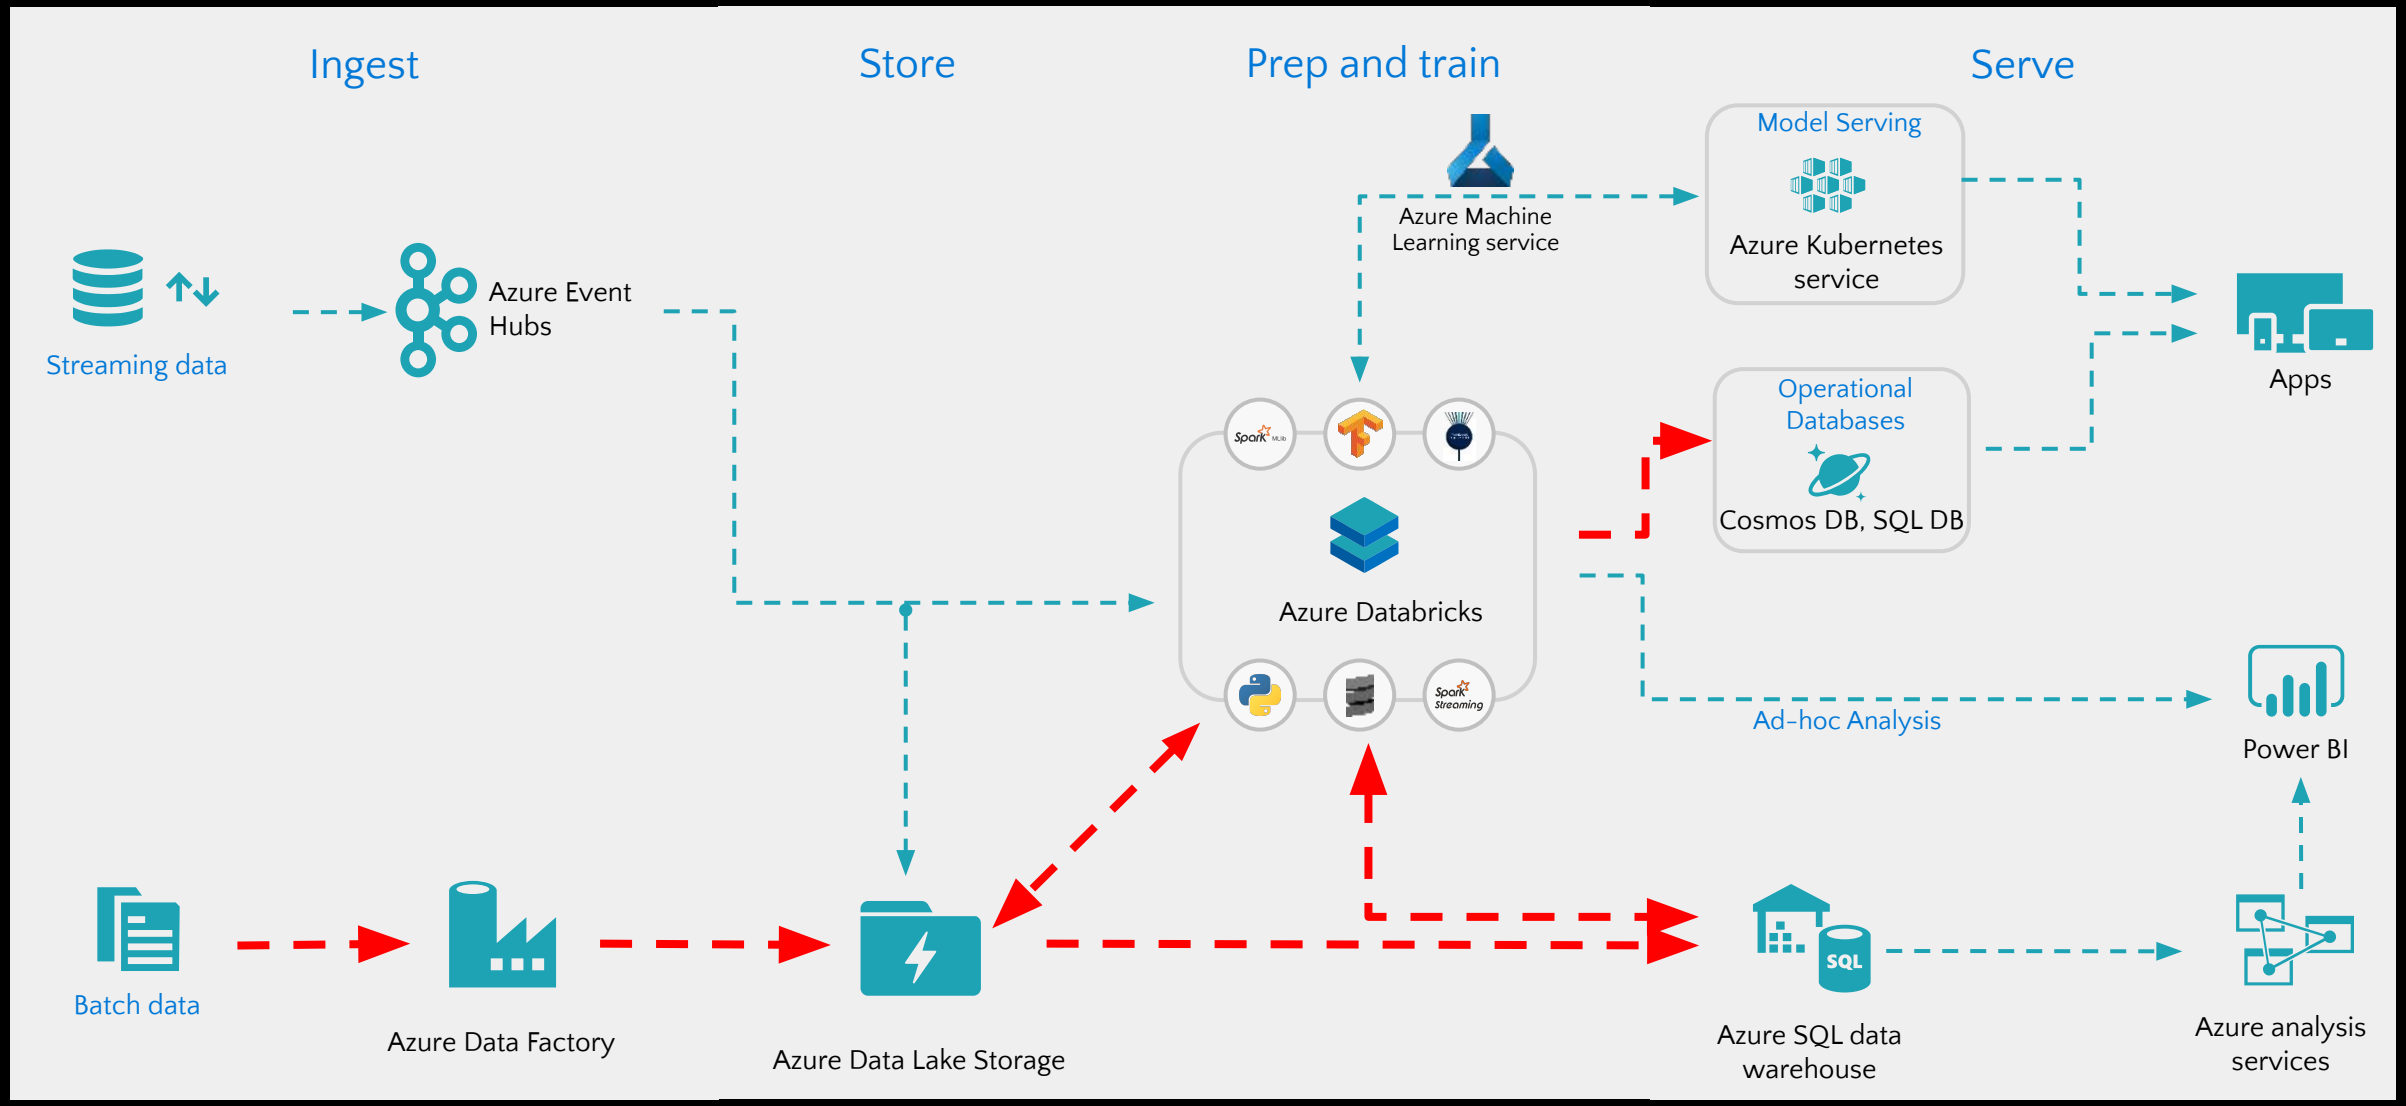

databricks

Recommended End-to-End Architecture - Stream ETL

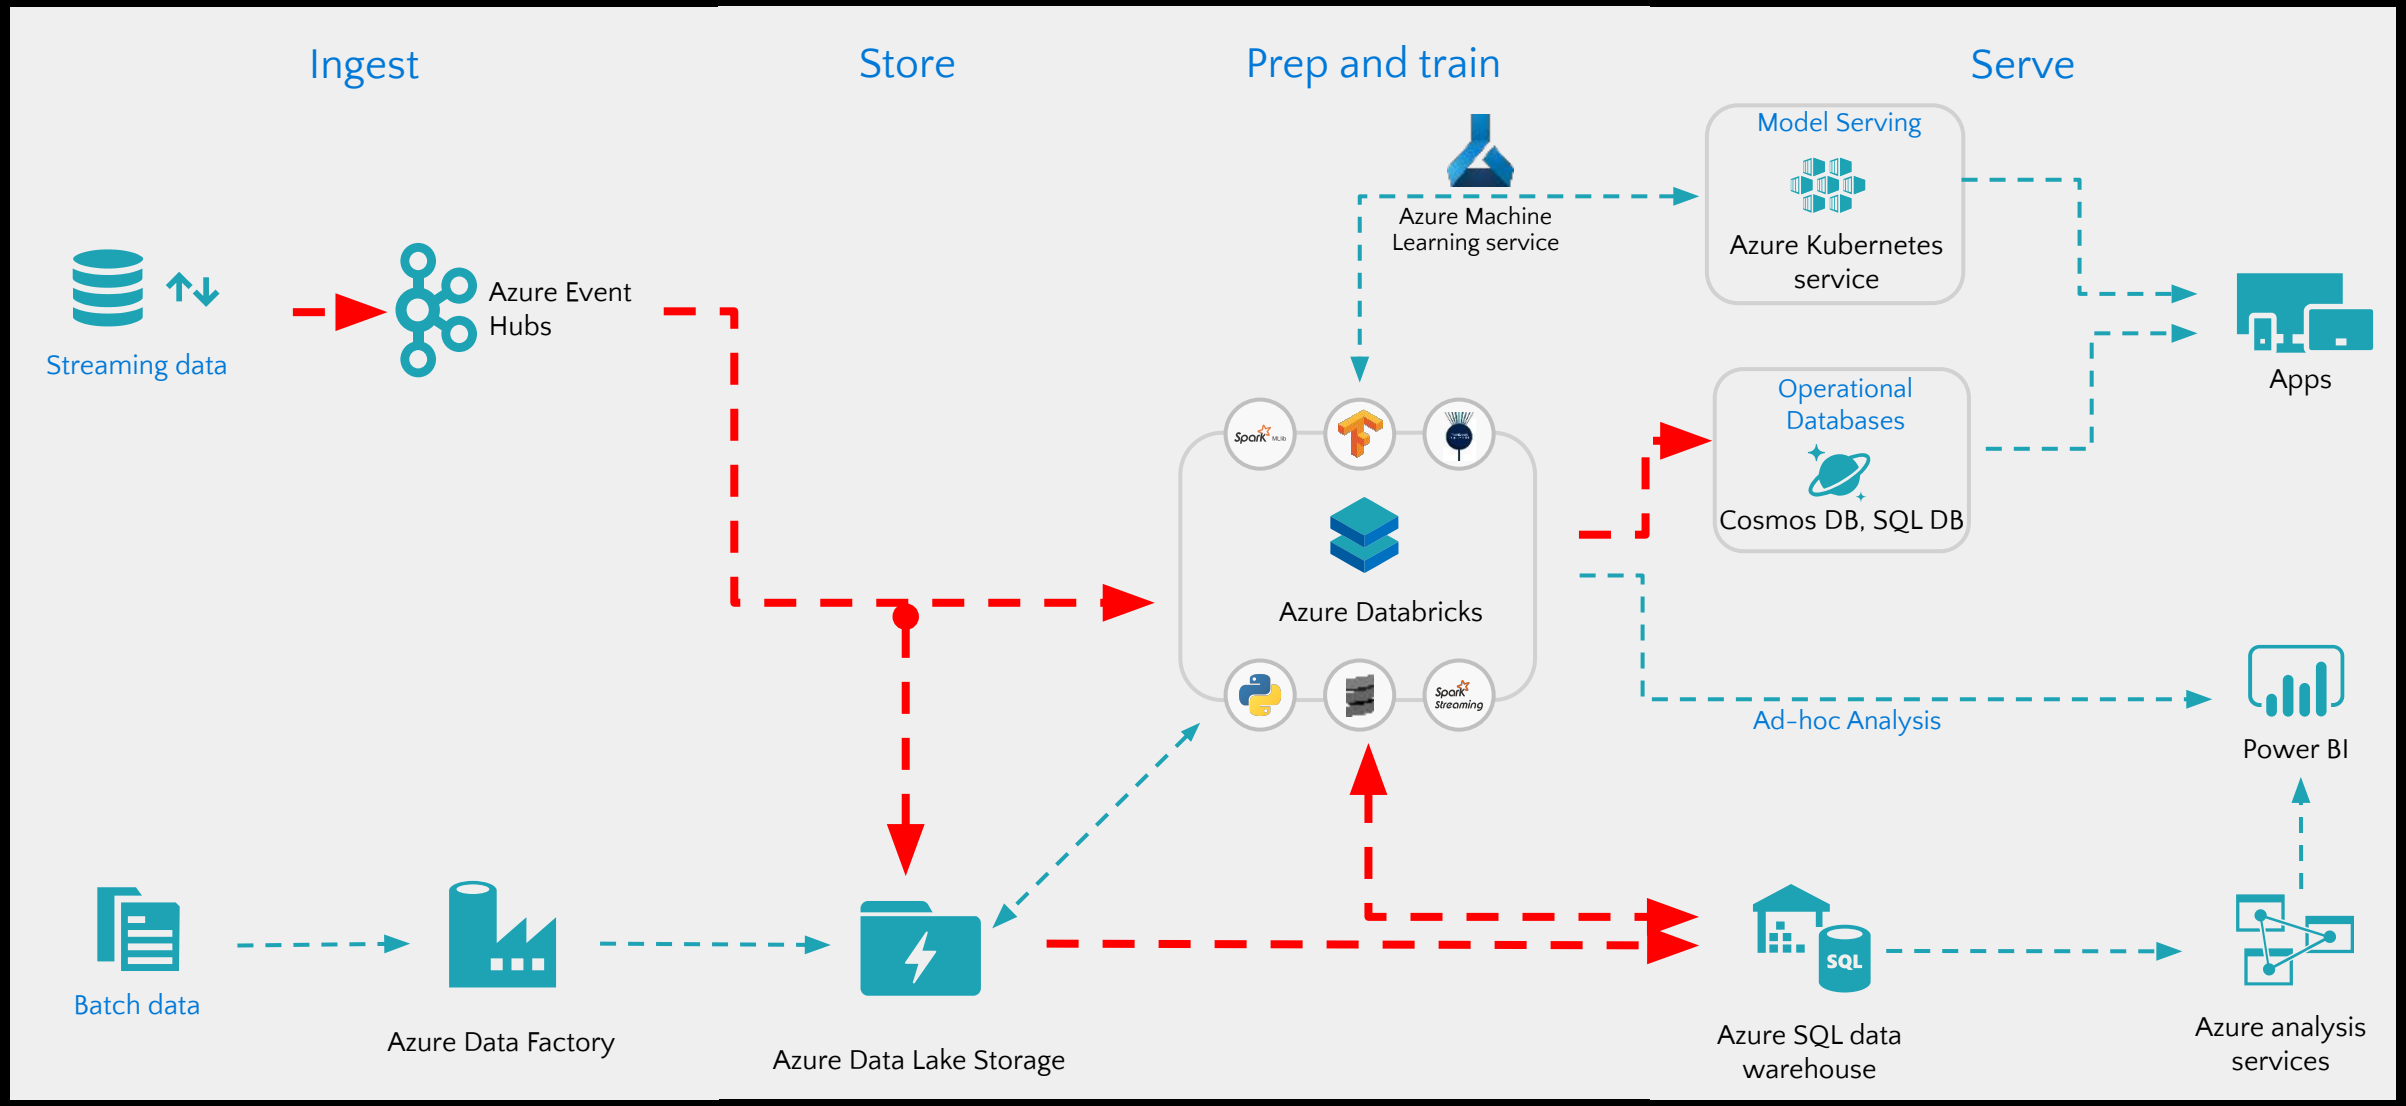

databricks

Recommended End-to-End Architecture - ML/Prediction

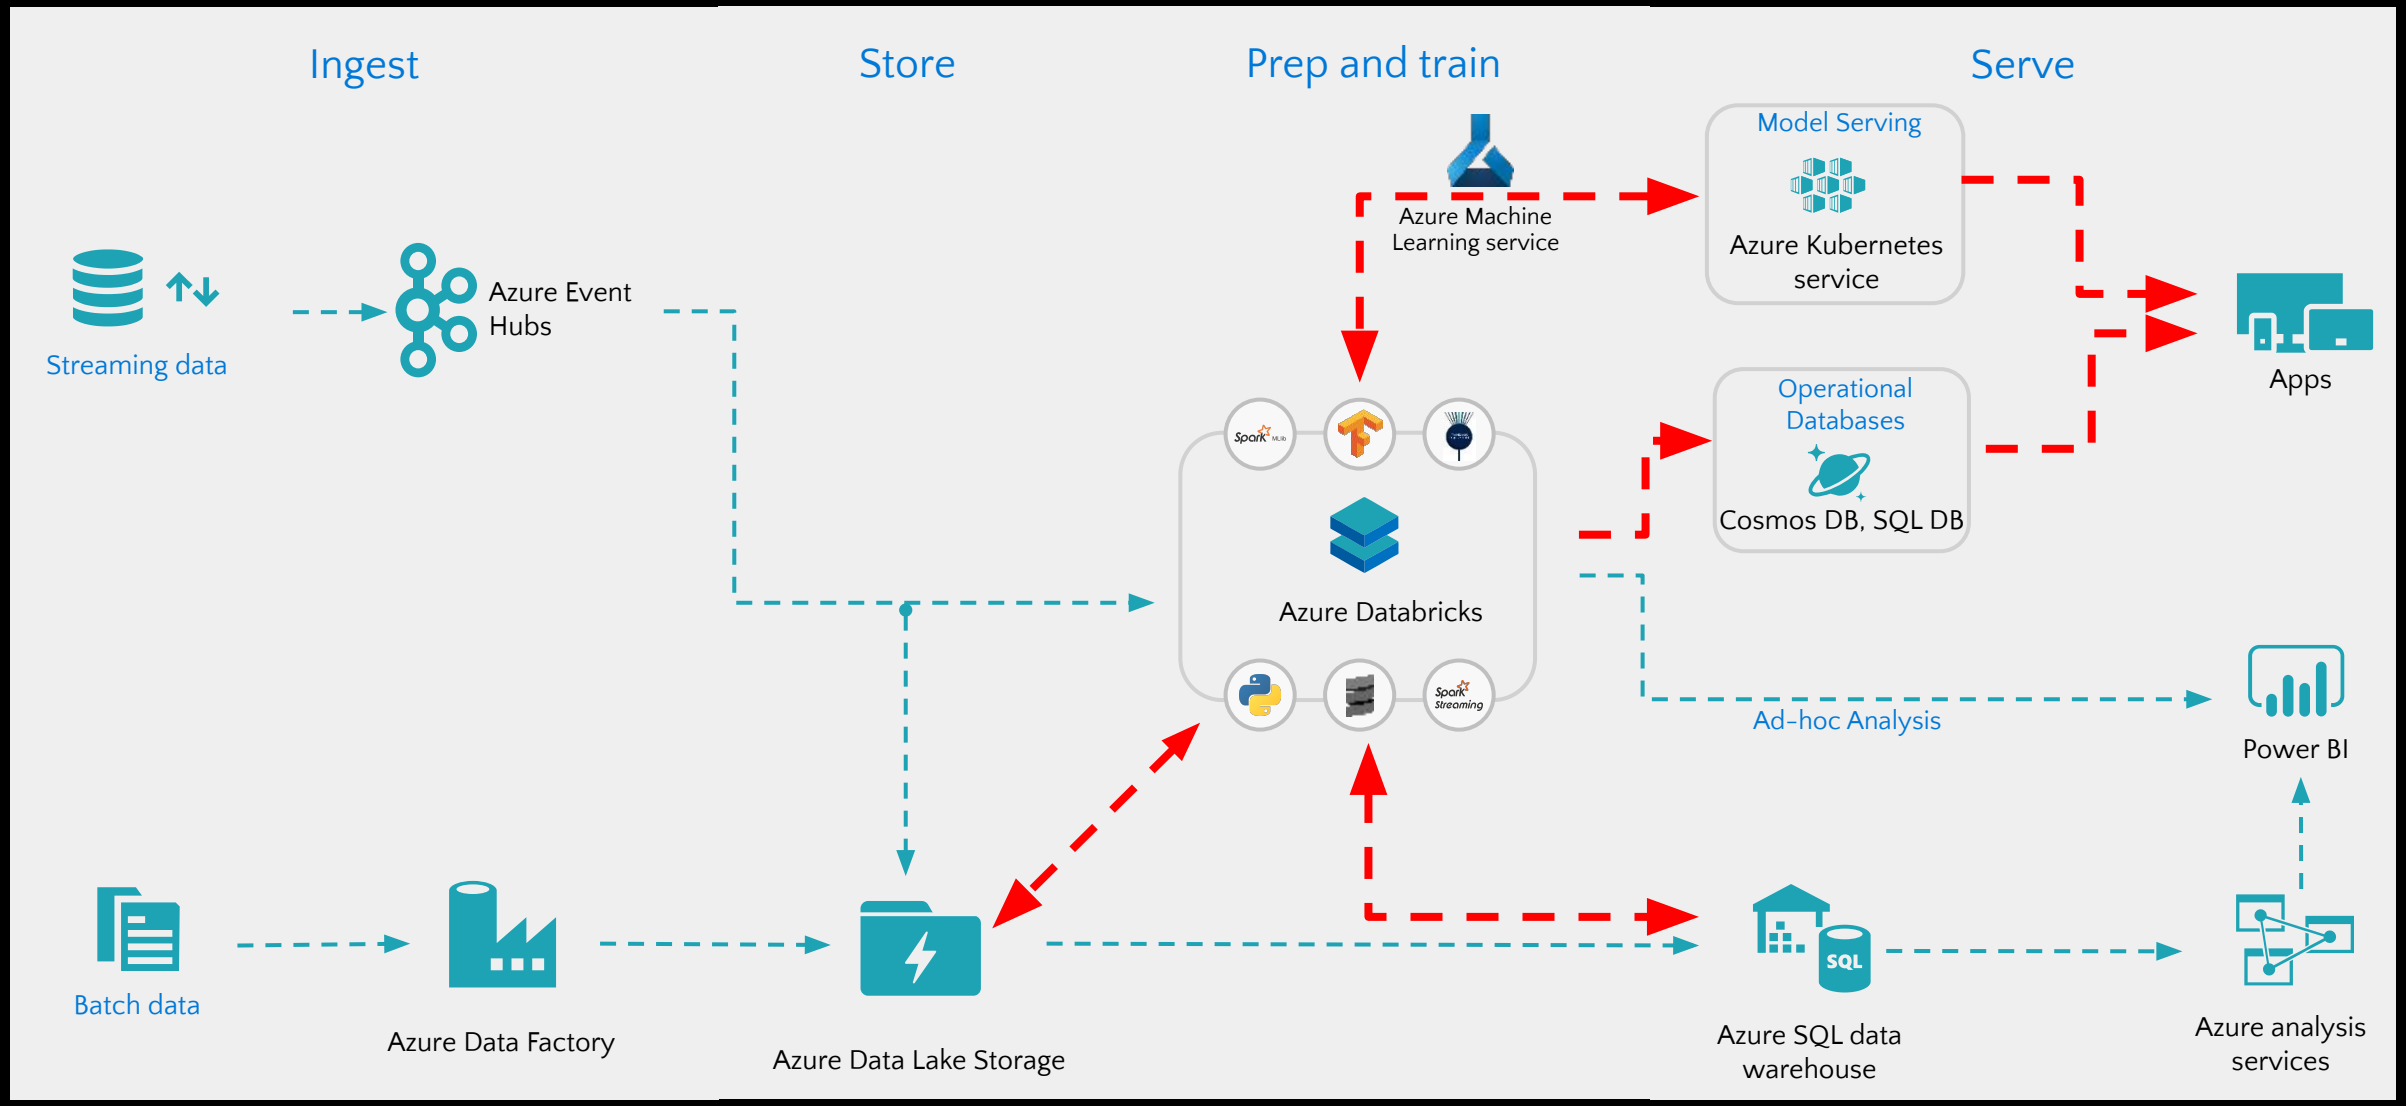

databricks

Recommended End-to-End Architecture - BI and Analysis

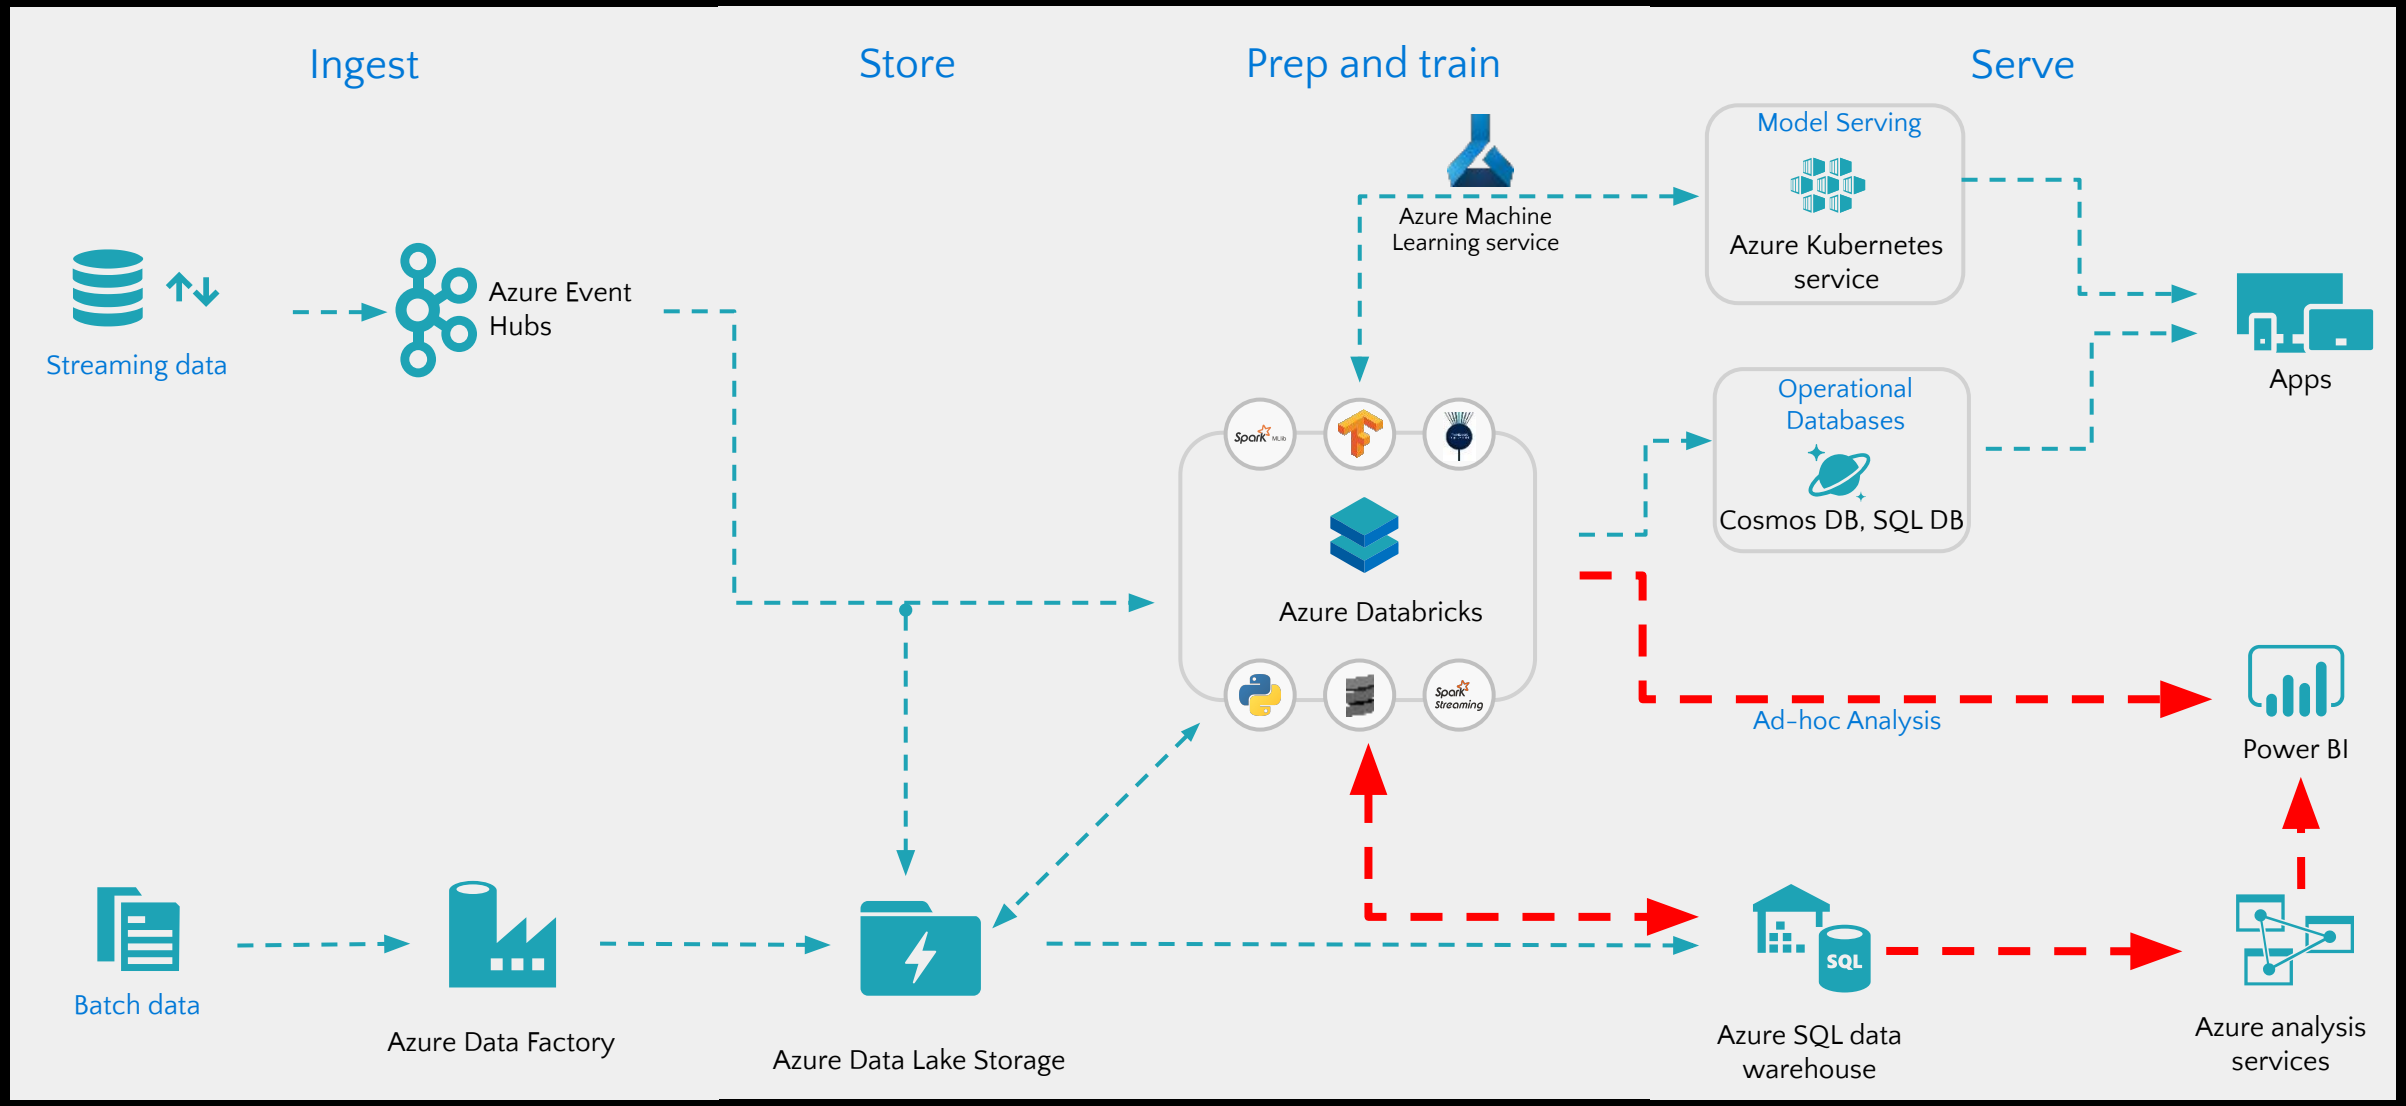

databricks

Recommended End-to-End Architecture

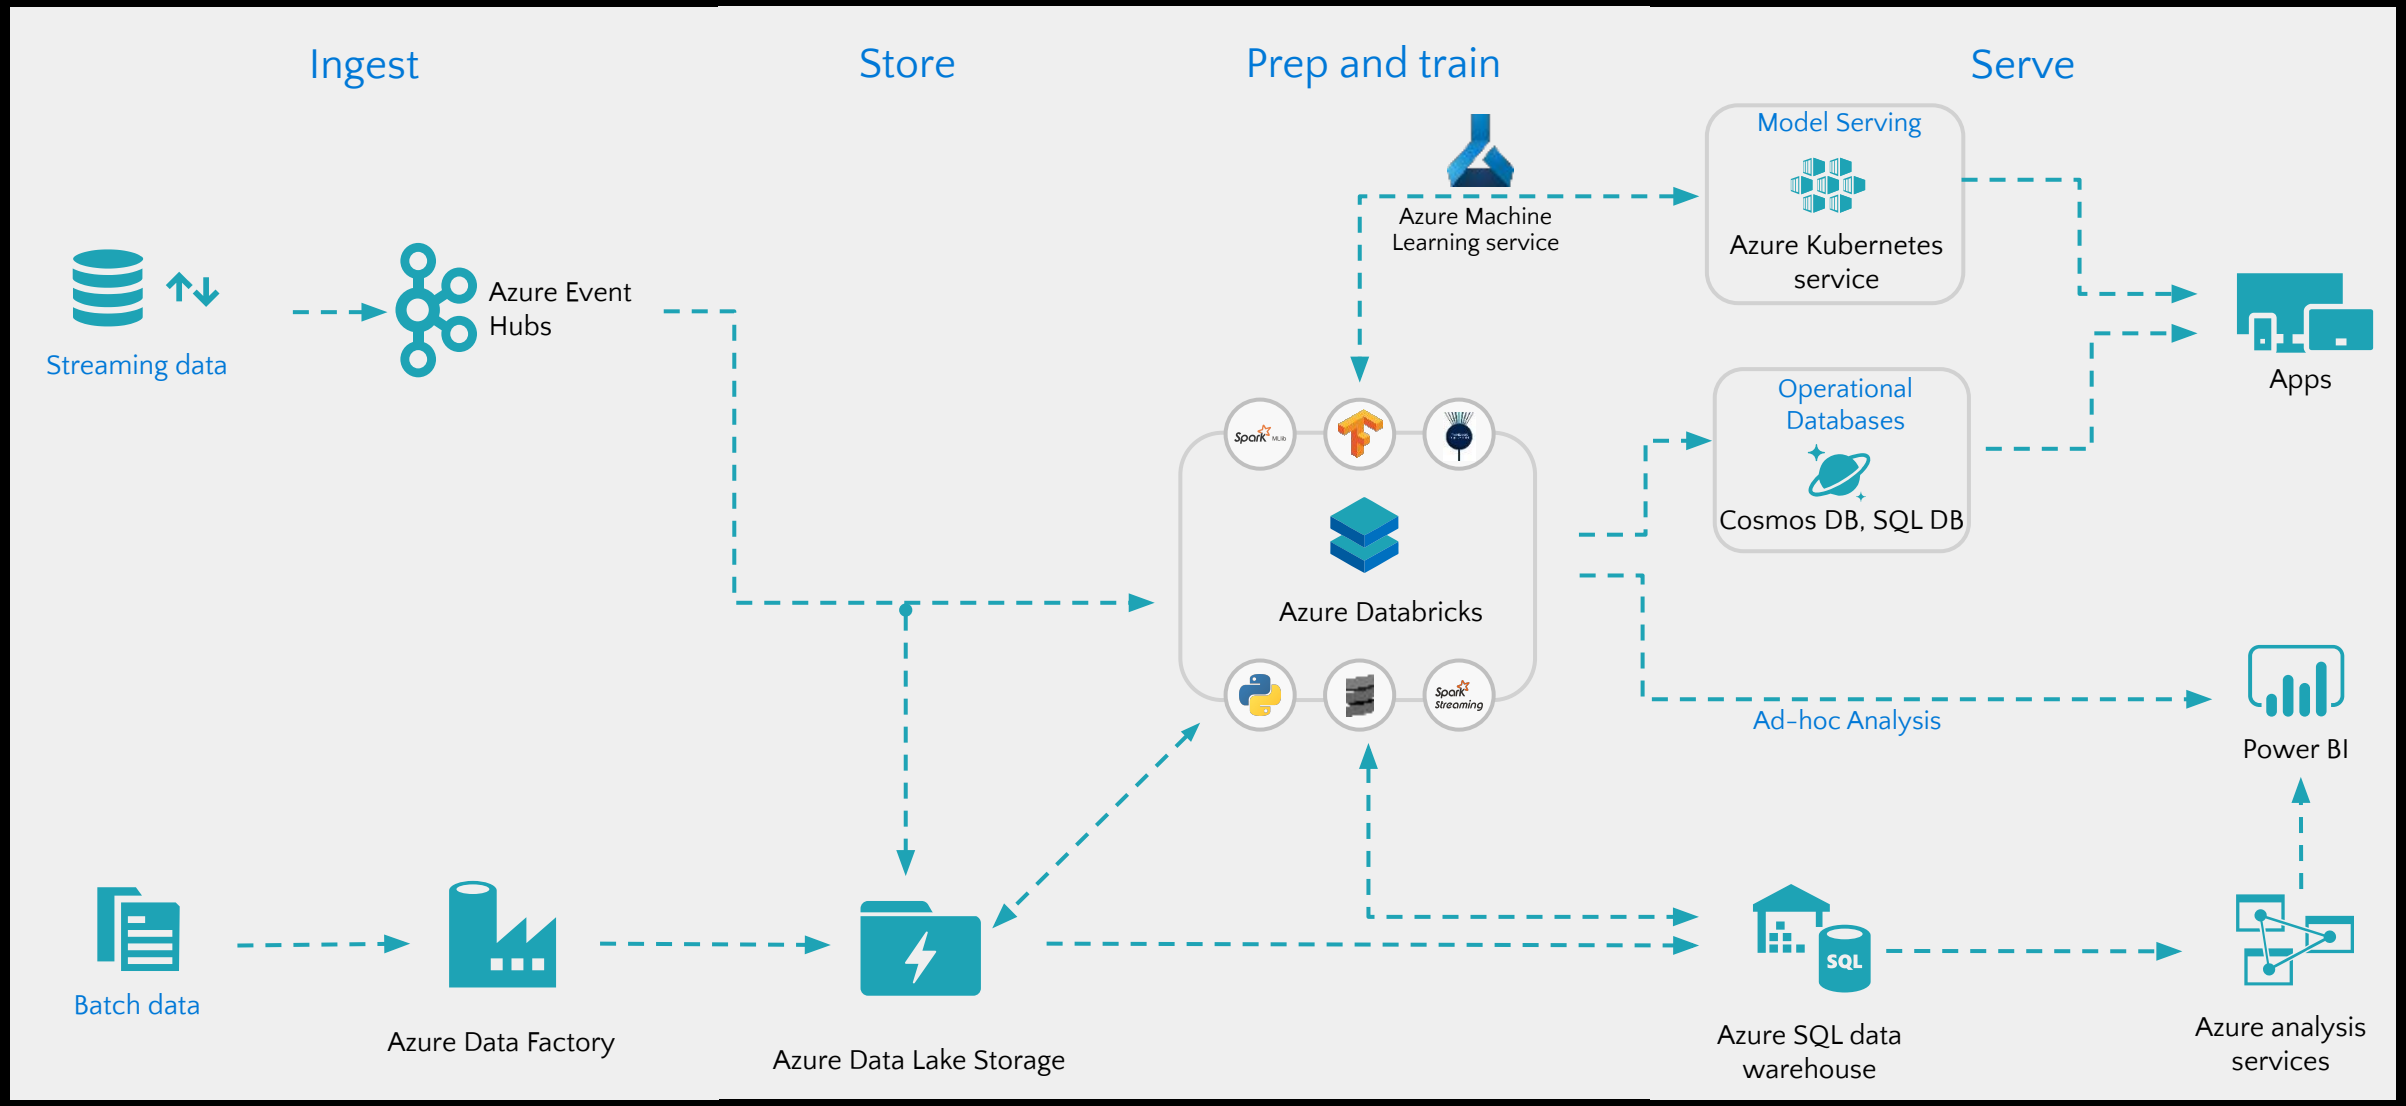

databricks

Azure Databricks Best Practices

Workspace Admin Best Practices

- Create different workspaces by different department / business team / data tier, and per environment (dev, qa, prod) - across relevant Azure subscriptions
- Define workspace level tags which propagate to initially provisioned resources in managed resource group (Tags could also propagate from parent resource group)
- Use ARM templates (search “databricks”) to have a more managed way of deploying the workspaces - whether via CLI, powershell or some SDK
- Create relevant groups of users - using Group REST API or by using AAD Group Sync with SCIM

Security Best Practices

- Do not store any production data on DBFS (use it only for toy / experimental datasets).
- Configure encryption-at-rest for Blob Storage and ADLS, preferably by using customer-managed keys in Azure Key Vault.
- Use Secrets with Azure Key Vault backend to obfuscate passwords and keys in notebooks.
- Prefer to use ADLS credential passthrough over Table ACLs (if possible).
- Configure access control for Databricks-native resources (clusters, notebooks, jobs etc.)
- Deploy workspace in your VNET to enable networking customizations.
- Configure Audit Logs to monitor the activity in a workspace.

Tools & Integration Best Practices

- Use Azure Data Factory to orchestrate pipelines / workflows (or something like Airflow).
- Connect your IDE or custom applications to Azure Databricks clusters using DB-Connect (Private Preview).
- Sync notebooks with Azure Devops for seamless version control.
- Use Databricks CLI for CI / CD from relevant enterprise tools/products, or to integrate with other systems like on-prem SCM or Library Repos etc.
- Use Library Utilities to install python libraries scoped at notebook level (cluster-scoped libraries may make more sense in certain cases).
- Use Init Scripts to do custom installs at cluster level.

Databricks Runtime Best Practices

- Use Delta wherever you can, to get the best performance and reliability for your big data workloads, and to create no-fuss multi-step data pipelines.
- Use Machine Learning Runtime for working with the latest ML/DL libraries (including HorovodRunner for distributed DL).
- Use DBIO Cache for accelerating reads from Blob Storage or ADLS.
- Use ABS-AQS connector for structured streaming when working with consistent rate of incoming files on Blob Storage.
- Turn on Databricks Advisor for automated tips on how to optimize workload processing.

HA and DR Best Practices

- Deploy Azure Databricks in two paired azure regions, ideally mapped to different control plane regions.
  - E.g. East US2 and West US2 will map to different control planes
  - Whereas West and North Europe will map to same control plane
- Use Azure Traffic Manager to load balance and distribute API requests between two deployments, when the platform is primarily being used in a backend non-interactive mode.
- Design to honor API and other limits of the platform.
  - Max API calls/ hr = 1500
  - Jobs per hour per workspace = 1000
  - Maximum concurrent Notebooks per cluster = 145

Cluster Best Practices

- Use autoscaling and auto-termination wherever applicable (e.g. auto-termination doesn't make sense if you need a cluster for data analysis by multiple users almost through the day, etc.).
- Use latest Databricks Runtime version to take advantage of latest performance & other optimizations (applicable in most cases, though not all).
- Use High-concurrency cluster mode for data analysis by a team of users via notebooks or a BI tool, or if you want to enforce data protection via Table ACLs or ADLS Passthrough.
- Use cluster tags for project / team based chargeback.

Cluster Best Practices Contd..

- Use Spark config tab if certain tuning would make sense for a specific workload (like config to use broadcast join).
- Use Event Log and Spark UI to see how different queries / workload executions perform, and what affect those have on a cluster's health.
- Configure Cluster Log Delivery
- Use Cluster ACLs to configure what each user or a group of users are allowed to do.
- Refer this blog by a customer, which more or less mentions what we've covered here. Rest is really workload dependent where it requires evidence-based tuning.

Appendix - Choosing the instance type

Different Azure Instance Types

Compute Optimized

- Fs
  - Haswell processor (Skylake not supported yet)
  - 1 core ~ 2GB RAM
  - SSD Storage: 1 core ~ 16GB
- H
  - High-performance
  - 1 core ~ 7GB RAM
  - SSD Storage: 1 core ~ 125GB

Memory Optimized

- DSv2
  - Haswell processor
  - 1 core ~ 7GB RAM
  - SSD Storage: 1 core ~ 14 GB
- ESv3
  - High-performance (Broadwell processor)
  - 1 core ~ 8GB RAM
  - SSD Storage: 1 core ~ 16GB

Storage Optimized

- L
  - 1 core ~ 8GB RAM
  - SSD Storage: 1 core ~ 170GB
  - Price : .156

General Purpose

- DSv2 and DSv3
  - DSv2 - 1 core ~ 3.5GB RAM
  - DSv3 - 1 core ~ 4GB RAM
  - SSD Storage:
    - DSv2 - 1 core ~ 7GB
    - DSv3 - 1 core ~ 8GB

Cluster Sizing Starting Points

Rules of Thumb

- Fewer big instances > more small instances
  - Reduce network shuffle; Databricks has 1 executor / machine
  - Applies to batch ETL mainly (for streaming, one could start with smaller instances depending on complexity of transformation)
  - Not set in stone, and reverse would make sense in many cases - so sizing exercise matters
- Size based on the number of tasks initially, tweak later
  - Run the job with a small cluster to get idea of # of tasks (use 2-3x tasks per core for base sizing)
- Choose based on workload (Probably start with F-series or DSv2):
  - ETL with full file scans and no data reuse - F / DSv2
  - ML workload with data caching - DSv2 / F
  - Data Analysis - L
  - Streaming - F

How do we tweak these?

Workload requires caching (like machine learning)

- Look at the Storage tab in Spark UI to see if the entirety of the training dataset is cached
  - Fully cached with room to spare -> less instances
  - Partially cached
    - Almost completely cached? -> Increase the cluster size
    - Not even close to cached -> Consider L series or DSv2 memory-optimized
      - Check to see if persist is MEMORY_ONLY, or MEMORY_AND_DISK
      - Spill to disk with SSD isn't so bad
- Still not good enough? Follow the steps in the next section

How do we tweak these?

ETL and Analytic Workloads

- Are we compute bound?
  - Check CPU Usage (Ganglia metrics to come to Azure Databricks soon)
  - Only way to make faster is more cores
- Are we network bound?
  - Check for high spikes before compute heavy steps
  - Use bigger/fewer machines to reduce the shuffle
  - Use an ssd backed instance for faster remote reads
- Are we spilling a ton?
  - Check Spark SQL tab for spill (pre-agg before shuffles are common to spill)
    - Use L-series
    - Or use more memory

Q&A